

IA-32 Architecture

Computer Organization and Assembly Languages

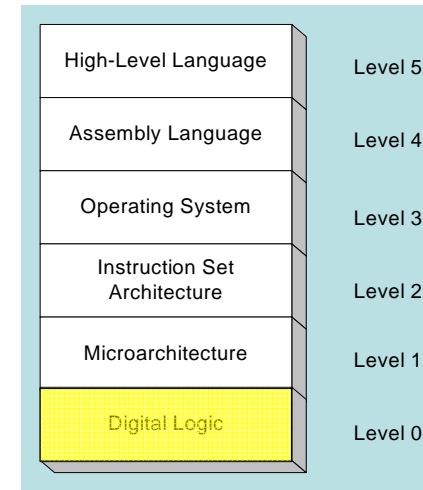
Yung-Yu Chuang

2005/10/6

with slides by Kip Irvine and Keith Van Rhein

Virtual machines


Abstractions for computers



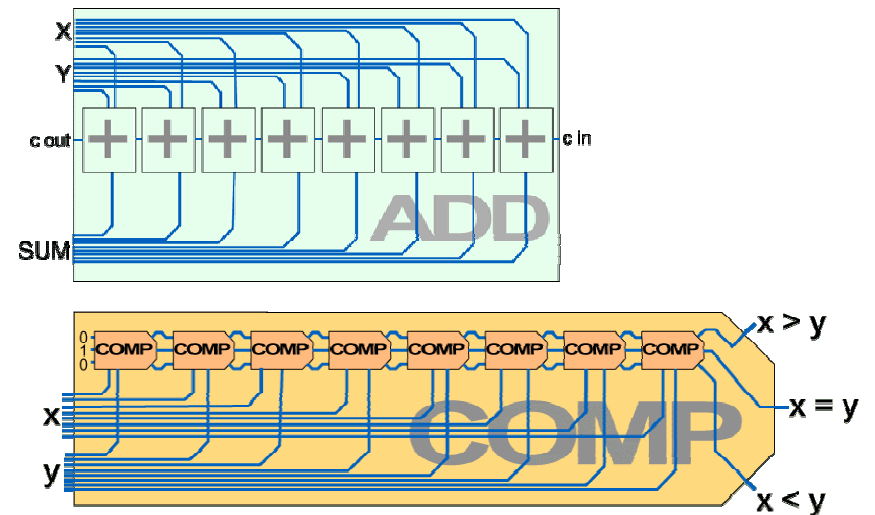
Truth tables

- Example: $(Y \wedge S) \vee (X \wedge \neg S)$


X	Y	S	$(Y \wedge S) \vee (X \wedge \neg S)$
F	F	F	F
F	T	F	F
T	F	F	T
T	T	F	T
F	F	T	F
F	T	T	T
T	F	T	F
T	T	T	T




Combinational logic



Sequential logic



Memory

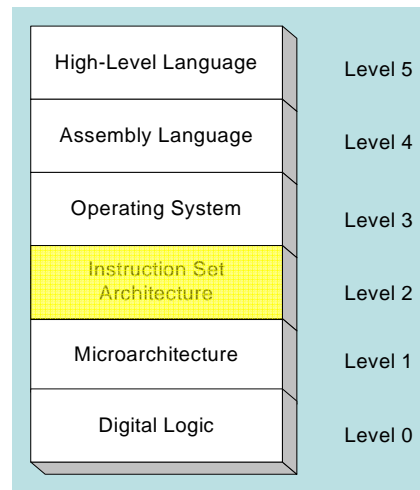


8K 8-bit memory

Virtual machines




Abstractions for computers



Instruction set

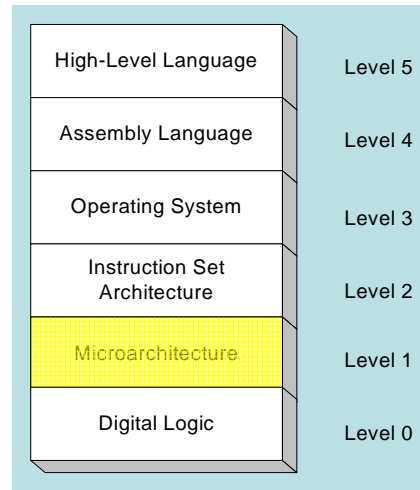


OPCODE	MNEMONIC	OPCODE	MNEMONIC
0	NOP	A	CMP addr
1	LDA addr	B	JG addr
2	STA addr	C	JE addr
3	ADD addr	D	JL addr
4	SUB addr		
5	IN port		
6	OUT port		
7	JMP addr		
8	JN addr		
9	HLT		



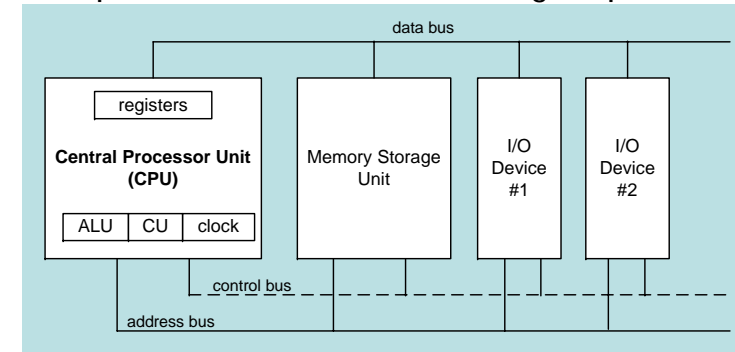
Virtual machines

Abstractions for computers



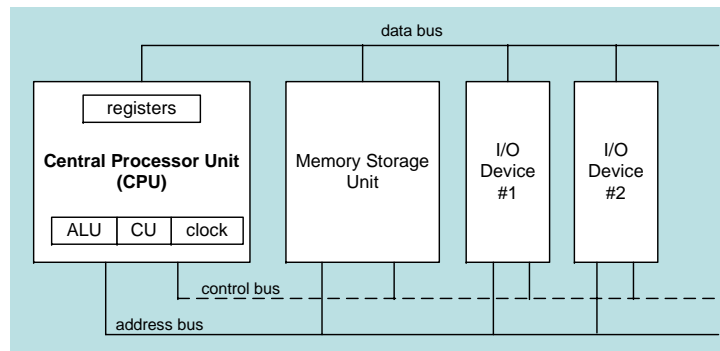
Basic microcomputer design

- clock synchronizes CPU operations
- control unit (CU) coordinates sequence of execution steps
- ALU performs arithmetic and logic operations




Basic microcomputer design

- The memory storage unit holds instructions and data for a running program
- A bus is a group of wires that transfer data from one part to another (data, address, control)



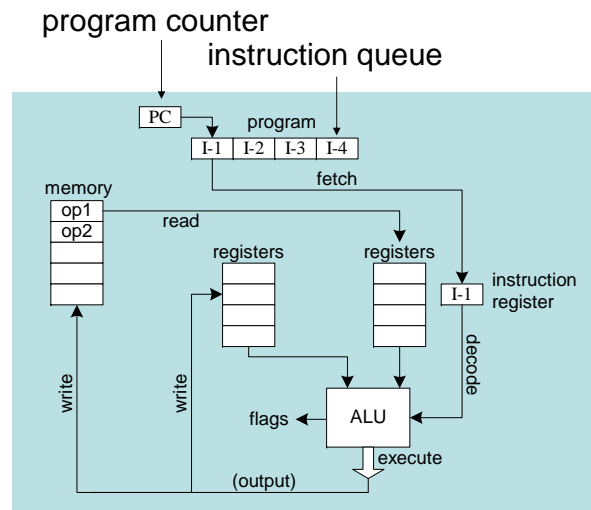
Clock

- synchronizes all CPU and BUS operations
- machine (clock) cycle measures time of a single operation
- clock is used to trigger events



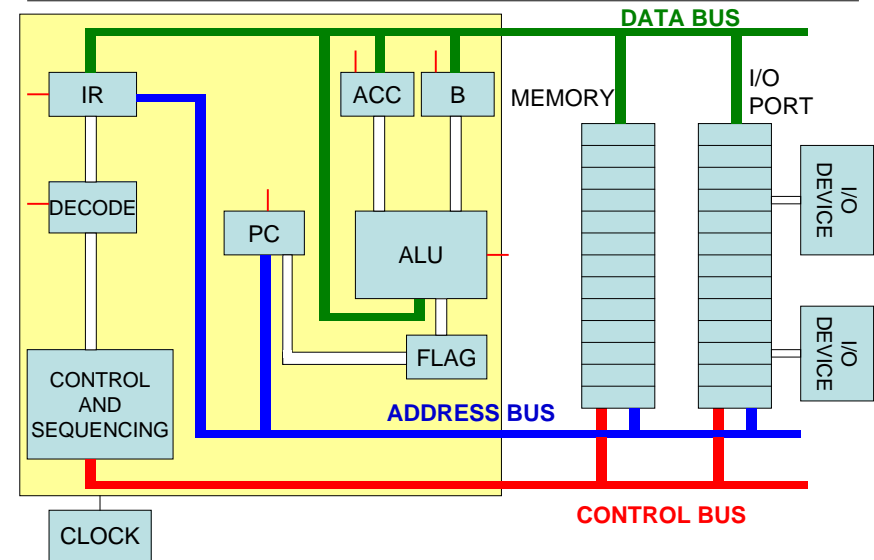
- Basic unit of time, 1GHz→clock cycle=1ns
- A instruction could take multiple cycles to complete, e.g. multiply in 8088 takes 50 cycles

Instruction execution cycle

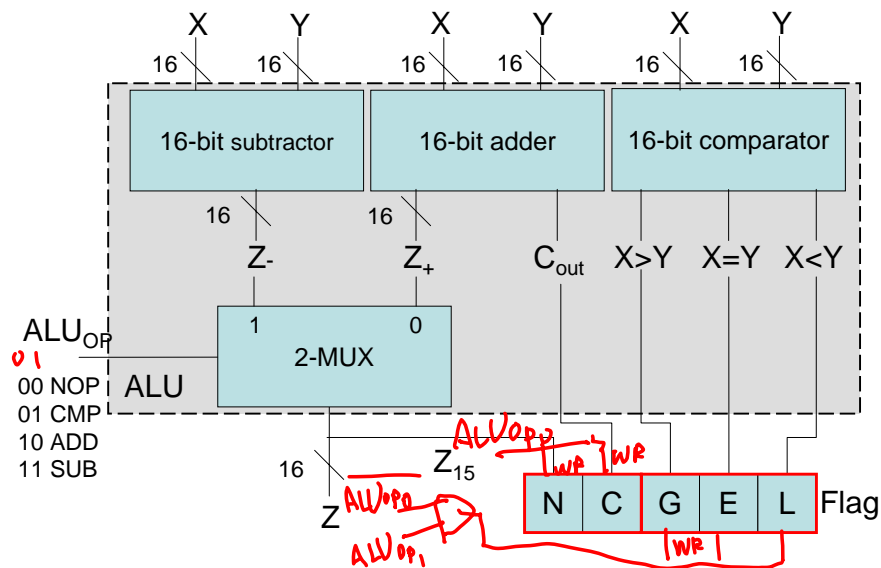


- Fetch
- Decode
- Fetch operands
- Execute
- Store output

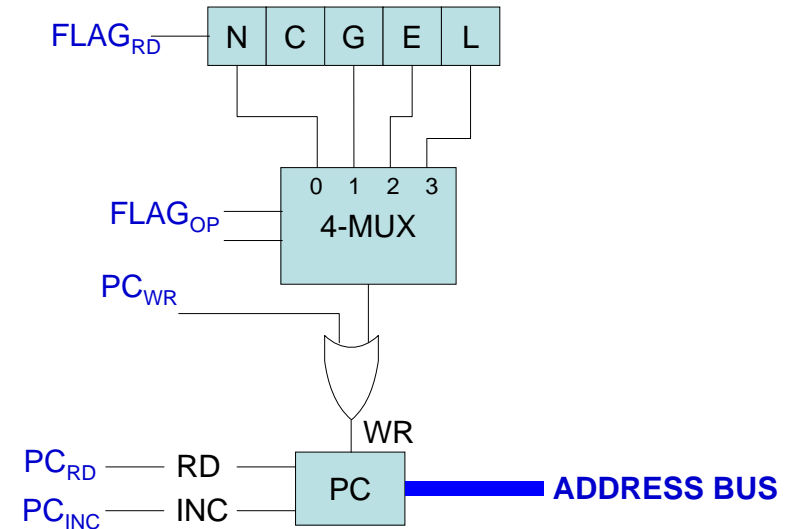
A simple microcomputer



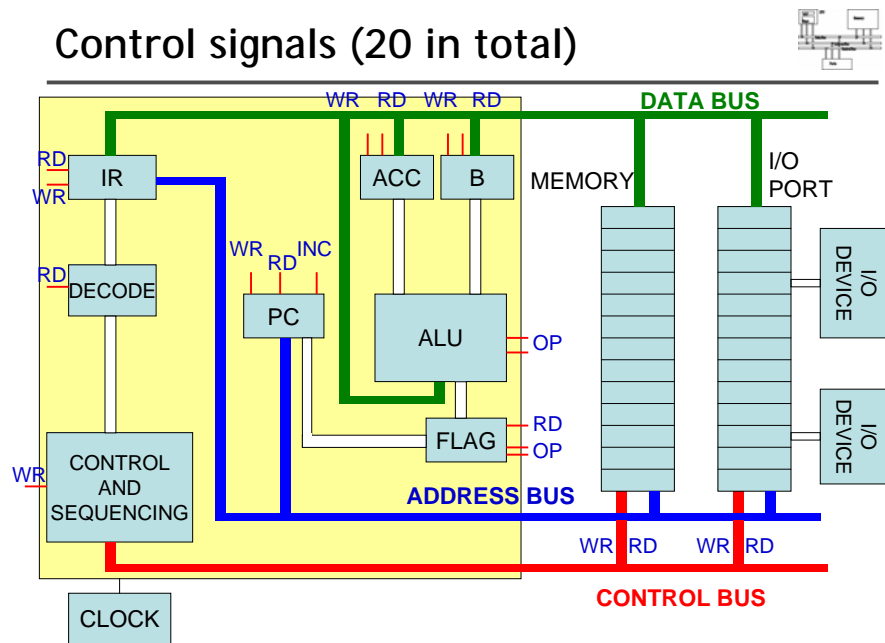
ALU and Flag



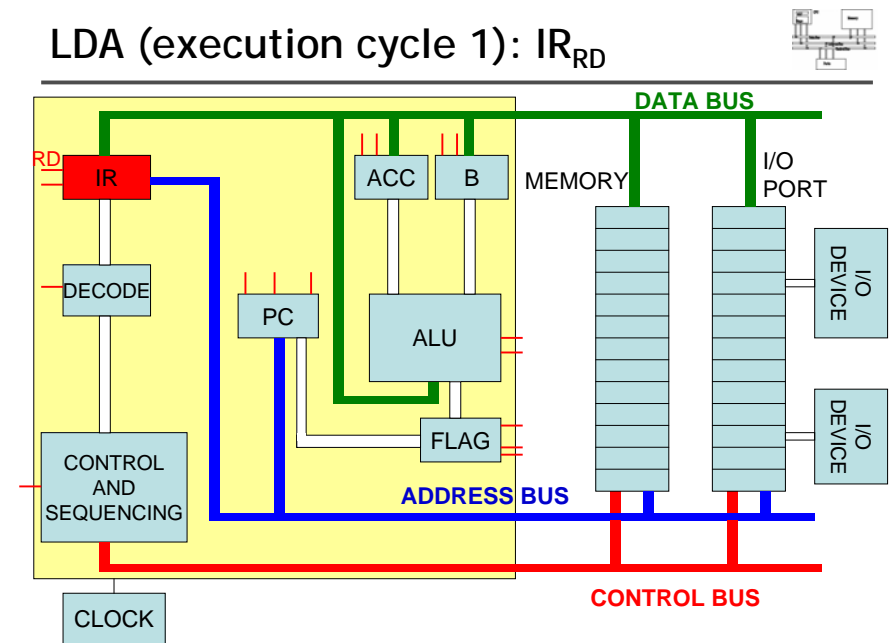
Flags



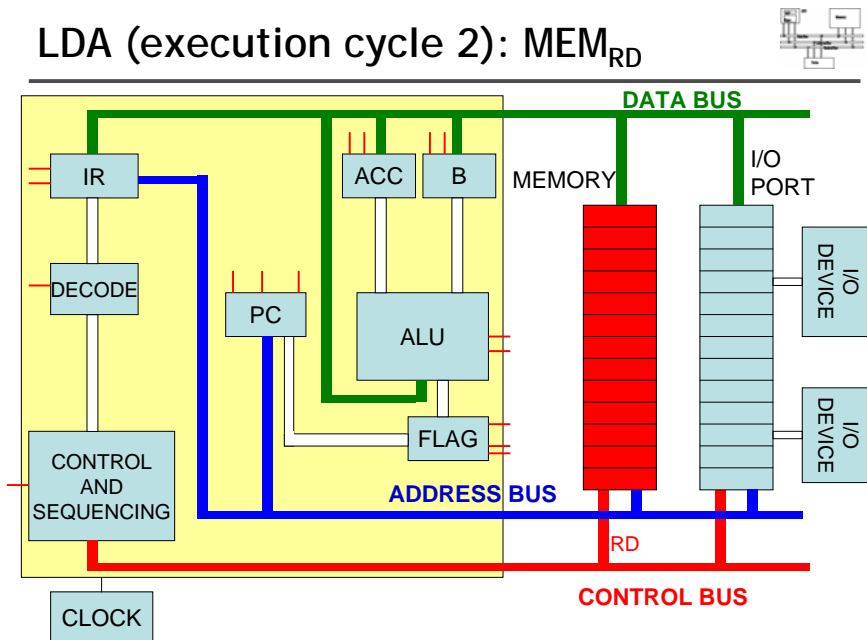
Control signals (20 in total)



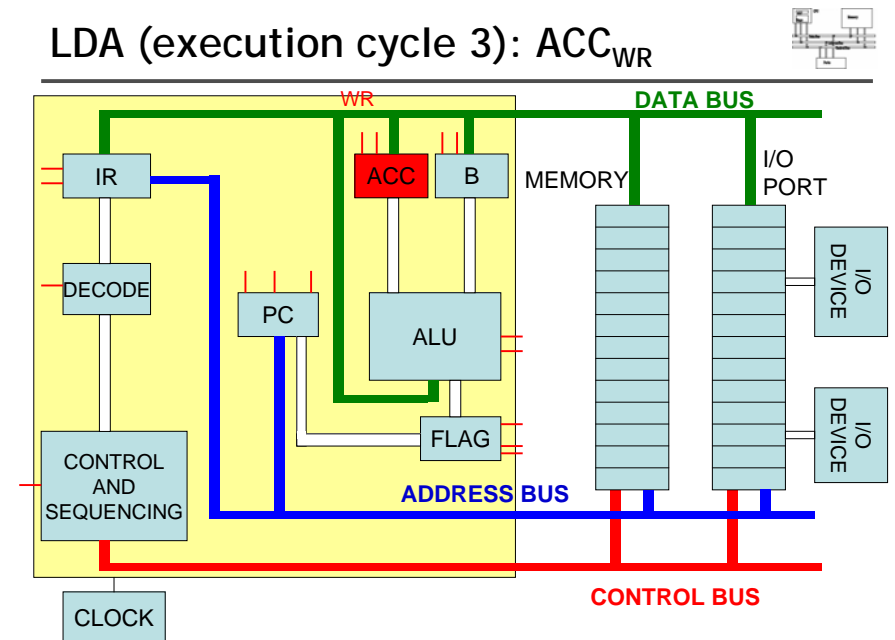
LDA (execution cycle 1): IR_{RD}




LDA (execution cycle 2): MEM_{RD}



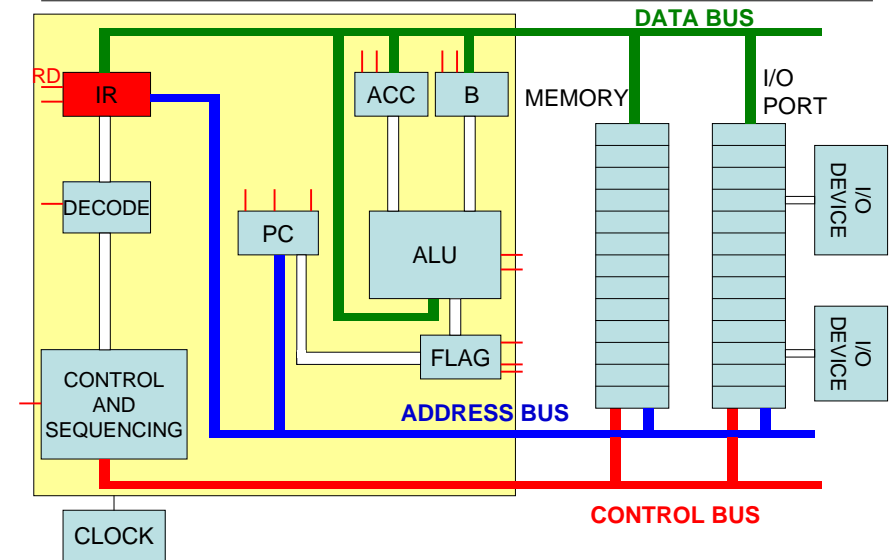
LDA (execution cycle 3): ACC_{WR}



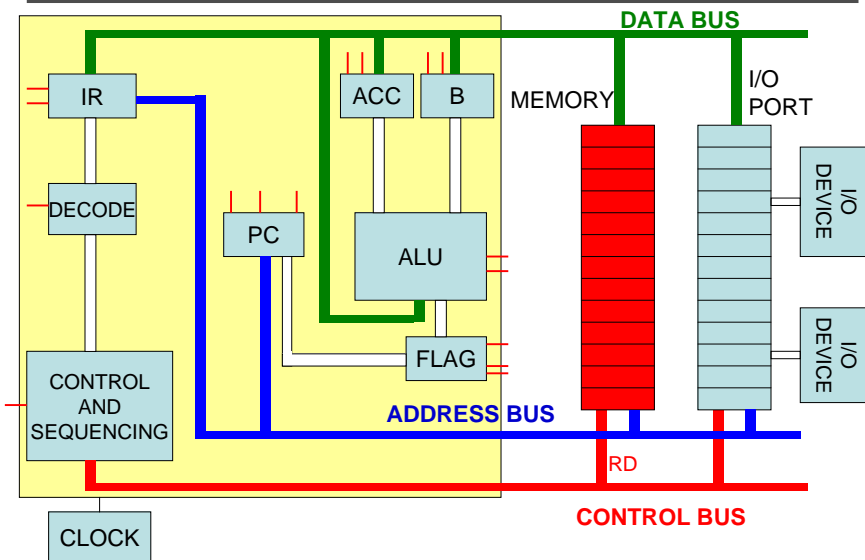
ALU and Flag



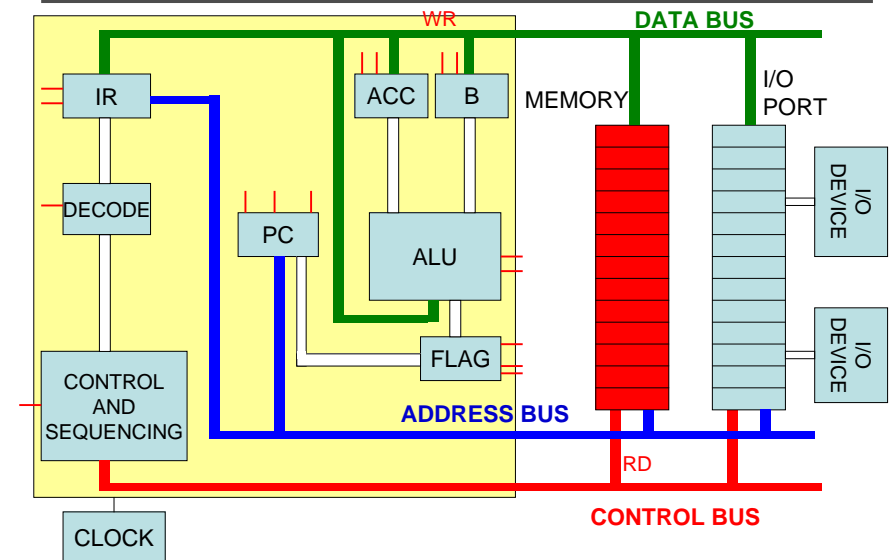
ADD (execution cycle 1): IR_{RD}



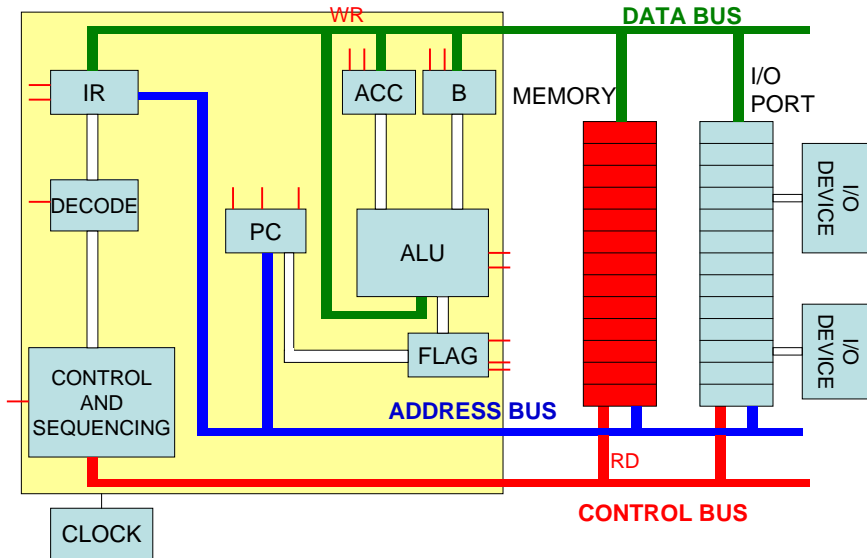
ADD (execution cycle 2): MEM_{RD}



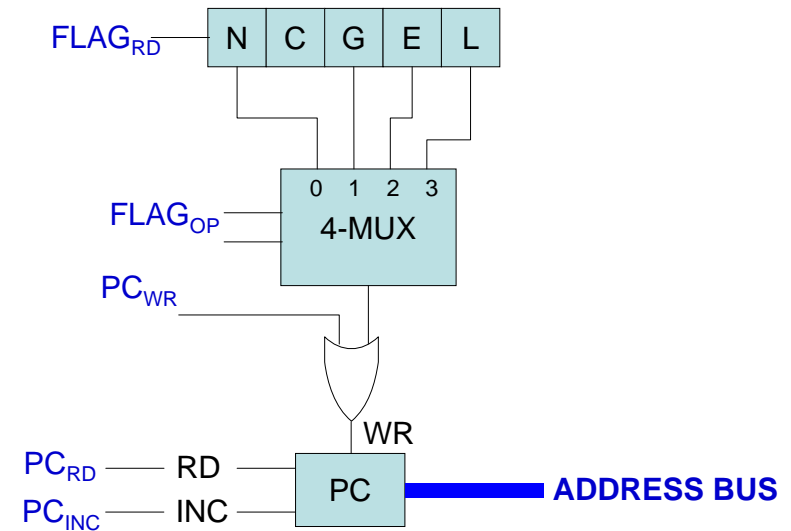
ADD (execution cycle 3): B_{WR}



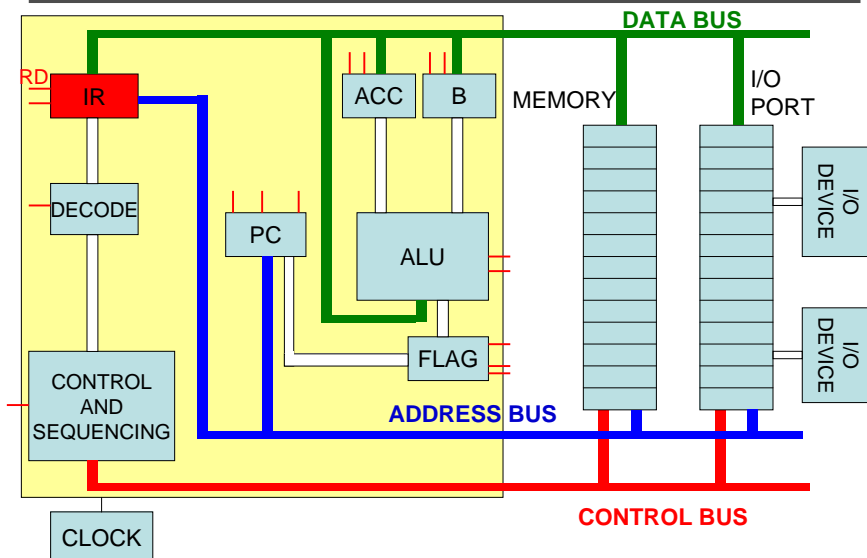
ADD (execution cycle 4): ALU_{10}, ACC_{WR}



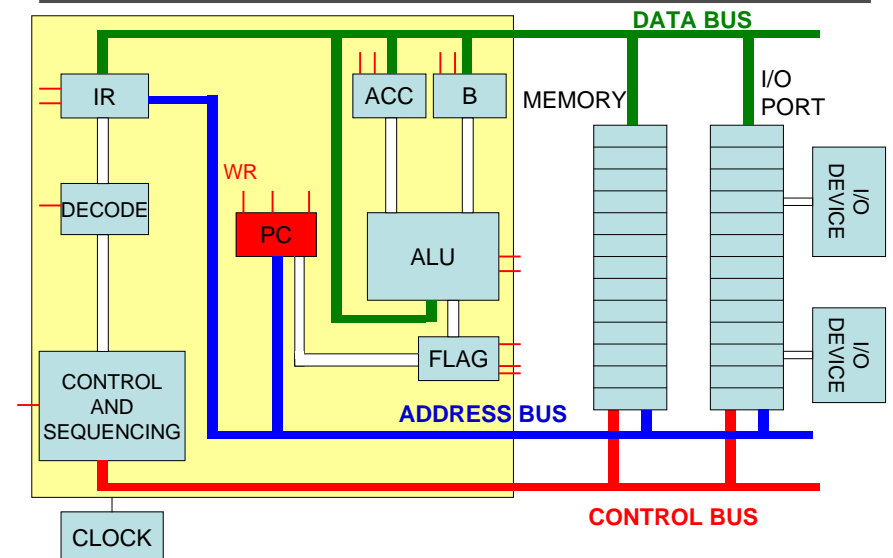
Flags



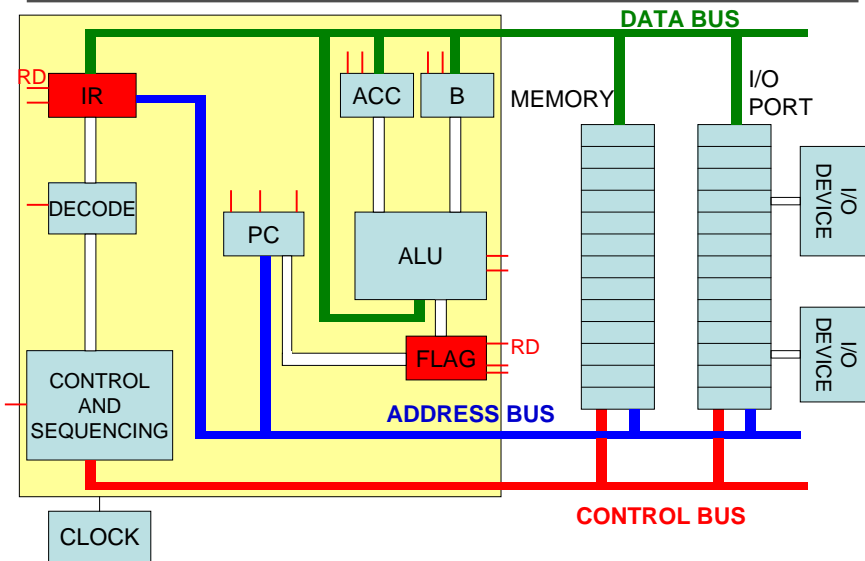
JMP (execution cycle 1): IR_{RD}



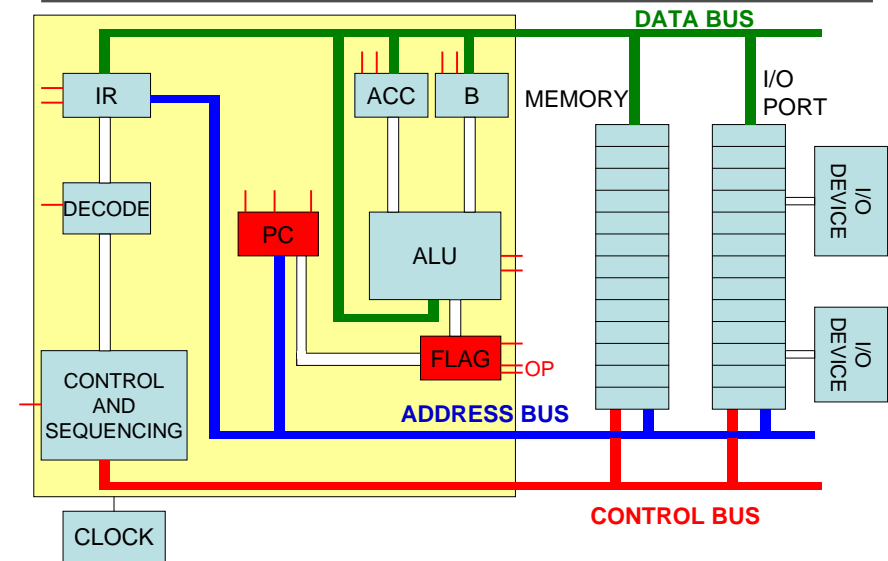
JMP (execution cycle 2): PC_{WR}



JG (execution cycle 1): $IR_{RD}, FLAG_{RD}$



JG (execution cycle 2): $FLAG_{01}$



Microcode sequence

LDA 510

PC_{RD}
MEM_{RD}
$IR_{WR} PC_{INC}$
IR_{RD}
$DECODE_{RD}$
μPC_{WR}
IR_{RD}
MEM_{RD}
ACC_{WR}
PC_{RD}
MEM_{RD}
$IR_{WR} PC_{INC}$
IR_{RD}
$DECODE_{RD}$
μPC_{WR}
IR_{RD}
PC_{WR}


JMP 10

Decoder

4-bit opcode

NOP	0	0000	μ code for LDA
LDA	1	0006	
STA	2	000F	
JMP	7		μ code for JMP

Control and sequencing unit



Control and sequencing unit

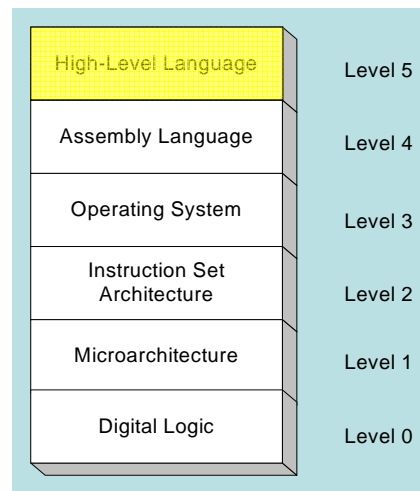


		PC _{RD}	MEM _{RD}	MEM _{WR}	IR _{WR}	PC _{INC}
NOP	0000	1	0	0	0	0	0....
fetch	0001	0	1	0	0	0	
	0002	0	0	0	1	1	
decode	0003	IR _{RD}					
	0004	DECODER _{RD}					
	0005	μ PC _{WR}					
LDAfetch							
decode							
exec	0006	IR _{RD}					
	0007	MEM _{RD}					
	0008	ACC _{WR}					
	000F						

Virtual machines



Abstractions for computers



X=min of X,Y,Z



```

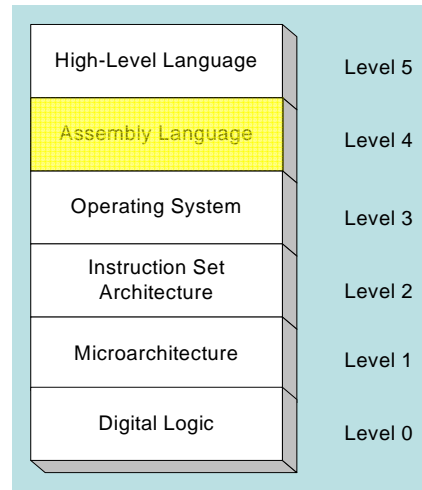
int X=7; Y=2; Z=9;
if (X>Y) then
    if (Y>Z) then
        X=Z;
    else
        X=Y;
    end
else
    if (X<Z) then
        X=Z;
    end ← else?
end

.DATA
X      007
Y      002
Z      009
.CODE
LDA    X
CMP    Y
JG     L1
CMP    Z
JL     L0
JMP    END
L0     LDA    Z
      STA    X
L1     LDA    Y
      CMP    Z
      JG     L2
      STA    X
      JMP    END
L2     LDA    Z
      STA    X
END     HLT
    
```


Virtual machines



Abstractions for computers



Memory layout



X=min of X,Y,Z

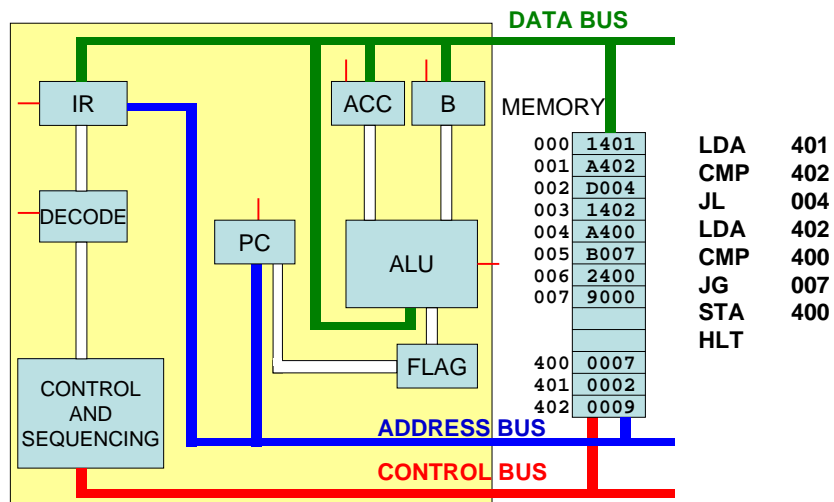


	.DATA		.DATA
	X 007		X 007
	Y 002		Y 002
	Z 009		Z 009
	.CODE		.CODE
	LDA X		LDA Y
	CMP Y		CMP Z
	JL L1		JL L1
	LDA Y		LDA Z
L1	CMP Z	L1	CMP X
	JL L2		JG END
	LDA Z		STA X
L2	STA X	END	HLT
	HLT		

X=min of X,Y,Z



	.DATA		
	X 007		X 400
	Y 002		Y 401
	Z 009		Z 402
	.CODE		
0	LDA Y	1401	
1	CMP Z	A402	
2	JL L1	D004	L1 4
3	LDA Z	1402	END 7
4	L1 CMP X	A400	
5	JG END	B007	
6	STA X	2400	
7	END HLT	9000	



Advanced architecture

Multi-stage pipeline



- Pipelining makes it possible for processor to execute instructions in parallel
- Instruction execution divided into discrete stages

Example of a non-pipelined processor. For example, 80386. Many wasted cycles.

Stages		S1	S2	S3	S4	S5	S6
Cycles	1	I-1					
	2		I-1				
	3			I-1			
	4				I-1		
	5					I-1	
	6						I-1
	7	I-2					
	8		I-2				
	9			I-2			
	10				I-2		
	11					I-2	
	12						I-2

Pipelined execution



- More efficient use of cycles, greater throughput of instructions: (80486 started to use pipelining)

Stages		S1	S2	S3	S4	S5	S6
Cycles	1	I-1					
	2	I-2	I-1				
	3		I-2	I-1			
	4			I-2	I-1		
	5				I-2	I-1	
	6					I-2	I-1
	7						I-2

For k stages and n instructions, the number of required cycles is:

$$k + (n - 1)$$

compared to $k \cdot n$

Wasted cycles (pipelined)



- When one of the stages requires two or more clock cycles, clock cycles are again wasted.

Cycles	Stages exe					
	S1	S2	S3	S4	S5	S6
	1	I-1				
	2	I-2	I-1			
	3	I-3	I-2	I-1		
	4		I-3	I-2	I-1	
	5			I-3	I-1	
	6				I-2	I-1
	7				I-2	
	8				I-3	I-2
	9				I-3	
	10					I-3
	11					

For k stages and n instructions, the number of required cycles is:

$$k + (2n - 1)$$

Superscalar



A superscalar processor has multiple execution pipelines. In the following, note that Stage S4 has left and right pipelines (u and v).

		Stages						
		S1	S2	S3	S4		S5	S6
					u	v		
Cycles	1	I-1						
	2	I-2	I-1					
	3	I-3	I-2	I-1				
	4	I-4	I-3	I-2	I-1			
	5		I-4	I-3	I-1	I-2		
	6			I-4	I-3	I-2	I-1	
	7				I-3	I-4	I-2	I-1
	8					I-4	I-3	I-2
	9						I-4	I-3
	10							I-4

For k states and n instructions, the number of required cycles is:

$$k + n$$

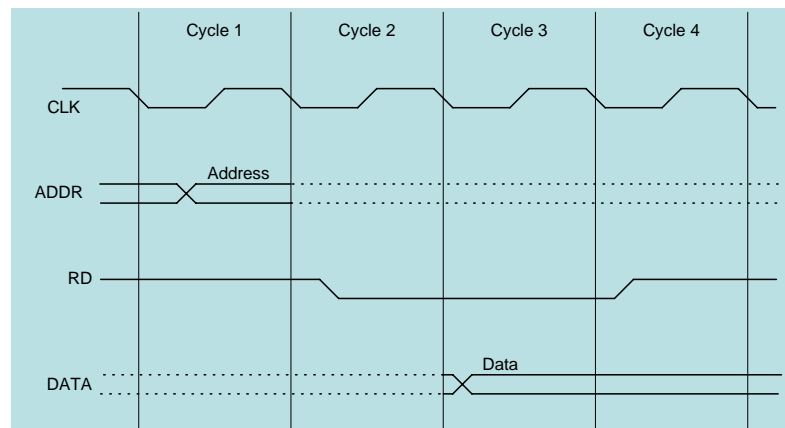
Pentium: 2 pipelines

Pentium Pro: 3


Reading from memory



- Multiple machine cycles are required when reading from memory, because it responds much more slowly than the CPU. The four steps are:

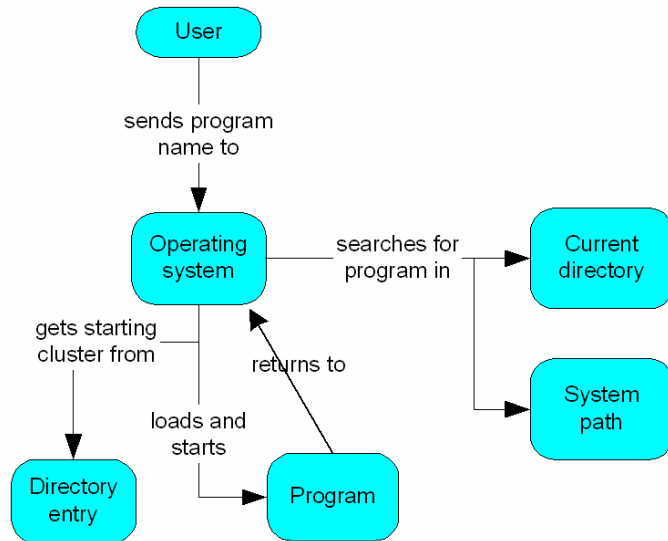


Cache memory



- High-speed expensive static RAM both inside and outside the CPU.
 - Level-1 cache: inside the CPU
 - Level-2 cache: outside the CPU
- Cache hit: when data to be read is already in cache memory
- Cache miss: when data to be read is not in cache memory. When? compulsory, capacity and conflict.
- Cache design: cache size, n-way, block size, replacement policy

How a program runs



Multitasking



- OS can run multiple programs at the same time.
- Multiple threads of execution within the same program.
- Scheduler utility assigns a given amount of CPU time to each running program.
- Rapid switching of tasks
 - gives illusion that all programs are running at once
 - the processor must support task switching
 - scheduling policy, round-robin, priority

IA-32 Architecture

IA-32 architecture



- From 386 to the latest 32-bit processor, P4
- From programmer's point of view, IA-32 has not changed substantially except the introduction of a set of high-performance instructions

Modes of operation



- Protected mode
 - native mode (Windows, Linux), full features, separate memory
- Virtual-8086 mode
 - hybrid of Protected
 - each program has its own 8086 computer
- Real-address mode
 - native MS-DOS
- System management mode
 - power management, system security, diagnostics

Addressable memory




- Protected mode
 - 4 GB
 - 32-bit address
- Real-address and Virtual-8086 modes
 - 1 MB space
 - 20-bit address

General-purpose registers




Named storage locations inside the CPU, optimized for speed.

32-bit General-Purpose Registers




16-bit Segment Registers



Accessing parts of registers



- Use 8-bit name, 16-bit name, or 32-bit name
- Applies to EAX, EBX, ECX, and EDX



32-bit	16-bit	8-bit (high)	8-bit (low)
EAX	AX	AH	AL
EBX	BX	BH	BL
ECX	CX	CH	CL
EDX	DX	DH	DL

Index and base registers



- Some registers have only a 16-bit name for their lower half. The 16-bit registers are usually used only in real-address mode.

32-bit	16-bit
ESI	SI
EDI	DI
EBP	BP
ESP	SP

Some specialized register uses (1 of 2)



- General-Purpose
 - EAX – accumulator (automatically used by division and multiplication)
 - ECX – loop counter
 - ESP – stack pointer (should never be used for arithmetic or data transfer)
 - ESI, EDI – index registers (used for high-speed memory transfer instructions)
 - EBP – extended frame pointer (stack)

Some specialized register uses (2 of 2)



- Segment
 - CS – code segment
 - DS – data segment
 - SS – stack segment
 - ES, FS, GS - additional segments
- EIP – instruction pointer
- EFLAGS
 - status and control flags
 - each flag is a single binary bit (*set* or *clear*)

Status flags




- Carry
 - unsigned arithmetic out of range
- Overflow
 - signed arithmetic out of range
- Sign
 - result is negative
- Zero
 - result is zero
- Auxiliary Carry
 - carry from bit 3 to bit 4
- Parity
 - sum of 1 bits is an even number

Floating-point, MMX, XMM registers



- Eight 80-bit floating-point data registers
 - ST(0), ST(1), . . . , ST(7)
 - arranged in a stack
 - used for all floating-point arithmetic
- Eight 64-bit MMX registers
- Eight 128-bit XMM registers for single-instruction multiple-data (SIMD) operations



IA-32 Memory Management

Real-address mode




- 1 MB RAM maximum addressable (20-bit address)
- Application programs can access any area of memory
- Single tasking
- Supported by MS-DOS operating system

Segmented memory



Segmented memory addressing: absolute (linear) address is a combination of a 16-bit segment value added to a 16-bit offset



Calculating linear addresses



- Given a segment address, multiply it by 16 (add a hexadecimal zero), and add it to the offset
- Example: convert 08F1:0100 to a linear address

Adjusted Segment value:	0 8 F 1 0
Add the offset:	0 1 0 0
Linear address:	0 9 0 1 0

- A typical program has three segments: code, data and stack. Segment registers CS, DS and SS are used to store them separately.

Example



What linear address corresponds to the segment/offset address 028F:0030?

$$028F0 + 0030 = 02920$$

Always use hexadecimal notation for addresses.

Example



What segment addresses correspond to the linear address 28F30h?

Many different segment-offset addresses can produce the linear address 28F30h. For example:

28F0:0030, 28F3:0000, 28B0:0430, . . .

Protected mode (1 of 2)



- 4 GB addressable RAM (32-bit address)
 - (00000000 to FFFFFFFFh)
- Each program assigned a memory partition which is protected from other programs
- Designed for multitasking
- Supported by Linux & MS-Windows

Protected mode (2 of 2)

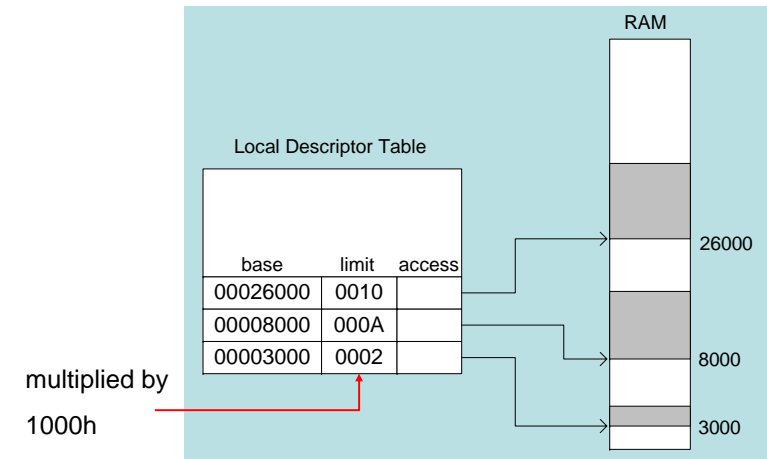


- Segment descriptor tables
- Program structure
 - code, data, and stack areas
 - CS, DS, SS segment descriptors
 - global descriptor table (GDT)
- MASM Programs use the Microsoft flat memory model

Multi-segment model




- Each program has a local descriptor table (LDT)
 - holds descriptor for each segment used by the program



Flat segmentation model



- All segments are mapped to the entire 32-bit physical address space, at least two, one for data and one for code
- global descriptor table (GDT)



Paging



- Virtual memory uses disk as part of the memory, thus allowing sum of all programs can be larger than physical memory
- Divides each segment into 4096-byte blocks called pages
- Page fault (supported directly by the CPU) – issued by CPU when a page must be loaded from disk
- Virtual memory manager (VMM) – OS utility that manages the loading and unloading of pages

Components of an IA-32 microcomputer


Components of an IA-32 Microcomputer

- Motherboard
- Video output
- Memory
- Input-output ports

Motherboard

- CPU socket
- External cache memory slots
- Main memory slots
- BIOS chips
- Sound synthesizer chip (optional)
- Video controller chip (optional)
- IDE, parallel, serial, USB, video, keyboard, joystick, network, and mouse connectors
- PCI bus connectors (expansion cards)

Intel D850MD motherboard



Video Output



- Video controller
 - on motherboard, or on expansion card
 - AGP (accelerated graphics port)
- Video memory (VRAM)
- Video CRT Display
 - uses raster scanning
 - horizontal retrace
 - vertical retrace
- Direct digital LCD monitors
 - no raster scanning required

Memory



- ROM
 - read-only memory
- EPROM
 - erasable programmable read-only memory
- Dynamic RAM (DRAM)
 - inexpensive; must be refreshed constantly
- Static RAM (SRAM)
 - expensive; used for cache memory; no refresh required
- Video RAM (VRAM)
 - dual ported; optimized for constant video refresh
- CMOS RAM
 - refreshed by a battery
 - system setup information

Input-output ports



- USB (universal serial bus)
 - intelligent high-speed connection to devices
 - up to 12 megabits/second
 - USB hub connects multiple devices
 - *enumeration*: computer queries devices
 - supports *hot* connections
- Parallel
 - short cable, high speed
 - common for printers
 - bidirectional, parallel data transfer
 - Intel 8255 controller chip

Input-output ports (cont)



- Serial
 - RS-232 serial port
 - one bit at a time
 - used for long cables and modems
 - 16550 UART (universal asynchronous receiver transmitter)
 - programmable in assembly language

Intel microprocessor history

Early Intel microprocessors



- Intel 8080
 - 64K addressable RAM
 - 8-bit registers
 - CP/M operating system
 - 5,6,8,10 MHz
 - 29K transistors
- Intel 8086/8088 (1978)
 - IBM-PC used 8088
 - 1 MB addressable RAM
 - 16-bit registers
 - 16-bit data bus (8-bit for 8088)
 - separate floating-point unit (8087)
 - used in low-cost microcontrollers now

The IBM-AT



- Intel 80286 (1982)
 - 16 MB addressable RAM
 - Protected memory
 - several times faster than 8086
 - introduced IDE bus architecture
 - 80287 floating point unit
 - Up to 20MHz
 - 134K transistors

Intel IA-32 Family



- Intel386 (1985)
 - 4 GB addressable RAM
 - 32-bit registers
 - paging (virtual memory)
 - Up to 33MHz
- Intel486 (1989)
 - instruction pipelining
 - Integrated FPU
 - 8K cache
- Pentium (1993)
 - Superscalar (two parallel pipelines)

Intel P6 Family



- Pentium Pro (1995)
 - advanced optimization techniques in microcode
 - More pipeline stages
 - On-board L2 cache
- Pentium II (1997)
 - MMX (multimedia) instruction set
 - Up to 450MHz
- Pentium III (1999)
 - SIMD (streaming extensions) instructions (SSE)
 - Up to 1+GHz
- Pentium 4 (2000)
 - NetBurst micro-architecture, tuned for multimedia
 - 3.8+GHz
- Pentium D (Dual core)

CISC and RISC



- CISC – complex instruction set
 - large instruction set
 - high-level operations (simpler for compiler?)
 - requires microcode interpreter (could take a long time)
 - examples: Intel 80x86 family
- RISC – reduced instruction set
 - small instruction set
 - simple, atomic instructions
 - directly executed by hardware very quickly
 - easier to incorporate advanced architecture design
 - examples:
 - ARM (Advanced RISC Machines)
 - DEC Alpha (now Compaq)