

Memory Hierarchy

Computer Organization and Assembly Languages

Yung-Yu Chuang

2006/01/05


with slides by CMU15-213

Announcement

- Grade for hw#4 is online
- Please DO submit homework if you haven't
- Please sign up a demo time on 1/16 or 1/17 at the following page
http://www.csie.ntu.edu.tw/~b90095/index.cgi/Assembly_Demo
- Hand in your report to TA at your demo time
- The length of report depends on your project type. It can be html, pdf, doc, ppt...


Reference

- Chapter 6 from "*Computer System: A Programmer's Perspective*"



Computer system model

- We assume memory is a linear array which holds both instruction and data, and CPU can access memory in a constant time.



SRAM vs DRAM




	Tran. per bit	Access time	Needs refresh?	Cost	Applications
SRAM	4 or 6	1X	No	100X	cache memories
DRAM	1	10X	Yes	1X	Main memories, frame buffers

The CPU-Memory gap



The gap widens between DRAM, disk, and CPU speeds.



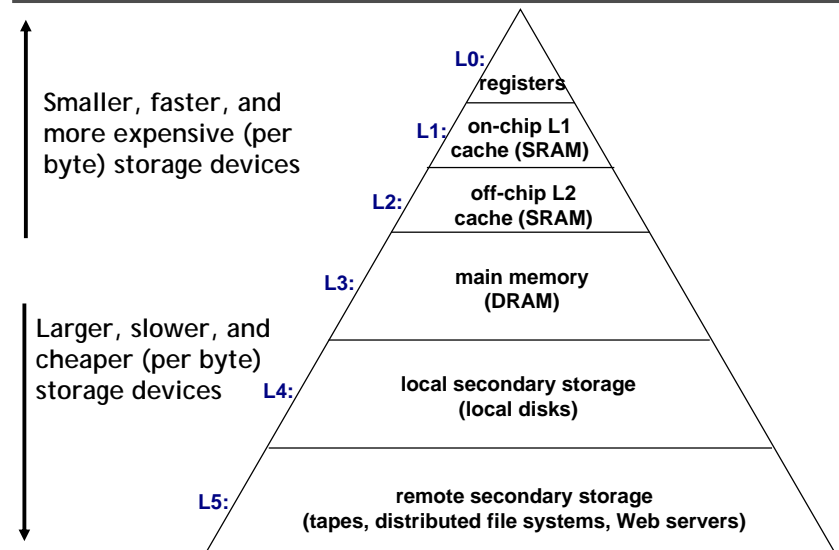
	register	cache	memory	disk
Access time (cycles)	1	1-10	50-100	20,000,000

Memory hierarchies



- Some fundamental and enduring properties of hardware and software:
 - Fast storage technologies cost more per byte, have less capacity, and require more power (heat!).
 - The gap between CPU and main memory speed is widening.
 - Well-written programs tend to exhibit good **locality**.
- They suggest an approach for organizing memory and storage systems known as a **memory hierarchy**.

Memory system in practice



Why it works?



- Most programs tend to access the storage at any particular level more frequently than the storage at the lower level.
- Locality: tend to access the same set of data items over and over again or tend to access sets of nearby data items.

Why learn it?



- A programmer needs to understand this because the memory hierarchy has a big impact on performance.
- You can optimize your program so that its data is more frequently stored in the higher level of the hierarchy.
- For example, the difference of running time for matrix multiplication could up to a factor of 6 even if the same amount of arithmetic instructions are performed.

Locality



- Principle of Locality: programs tend to reuse data and instructions near those they have used recently, or that were recently referenced themselves.
 - **Temporal locality**: recently referenced items are likely to be referenced in the near future.
 - **Spatial locality**: items with nearby addresses tend to be referenced close together in time.
- In general, *programs with good locality run faster than programs with poor locality*
- Locality is the reason why cache and virtual memory are designed in architecture and operating system. Another example is web browser caches recently visited webpages.

Locality example



```
sum = 0;
for (i = 0; i < n; i++)
    sum += a[i];
return sum;
```


- **Data**
 - Reference array elements in succession (stride-1 reference pattern): **Spatial locality**
 - Reference sum each iteration: **Temporal locality**
- **Instructions**
 - Reference instructions in sequence: **Spatial locality**
 - Cycle through loop repeatedly: **Temporal locality**

Locality example

- Being able to look at code and get a qualitative sense of its locality is important. Does this function have good locality?

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
} stride-1 reference pattern
```




Locality example

- Does this function have good locality?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
} stride-N reference pattern
```



Locality example

```
typedef struct {
    float v[3];
    float a[3];
} point ;
point p[N];
```

```
for (i=0; i<n; i++) {
    for (j=0; j<3; j++) {
        p[i].v[j]=0;
        p[i].a[j]=0;
    }
}
```

```
for (i=0; i<n; i++) {
    for (j=0; j<3; j++)
        p[i].v[j]=0;
    for (j=0; j<3; j++)
        p[i].a[j]=0;
}
```

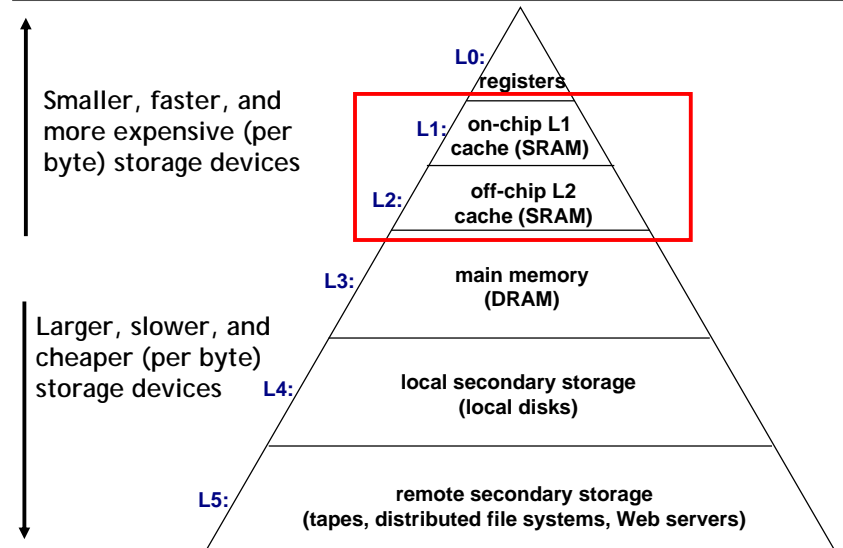
(A)

```
for (j=0; j<3; j++) {
    for (i=0; i<n; i++)
        p[i].v[j]=0;
    for (i=0; i<n; i++)
        p[i].a[j]=0;
}
```

(C)

(B)

Memory hierarchies

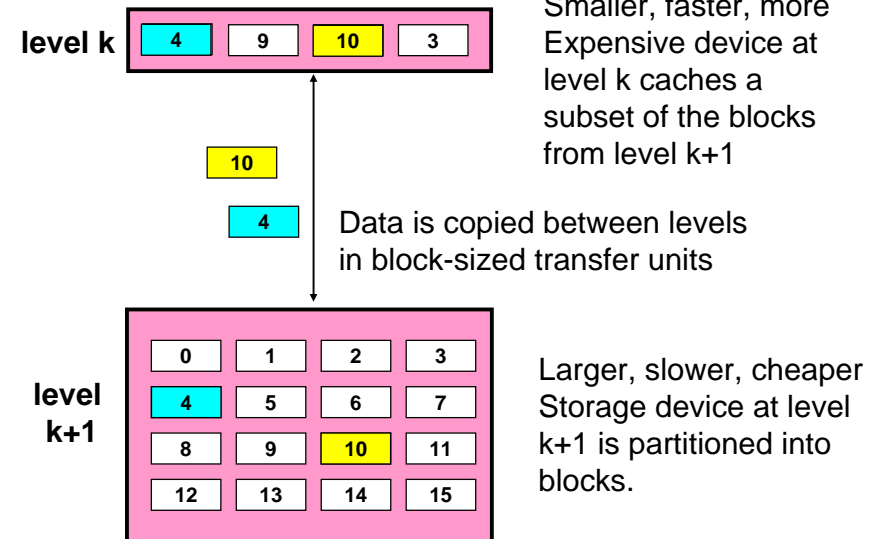


Caches

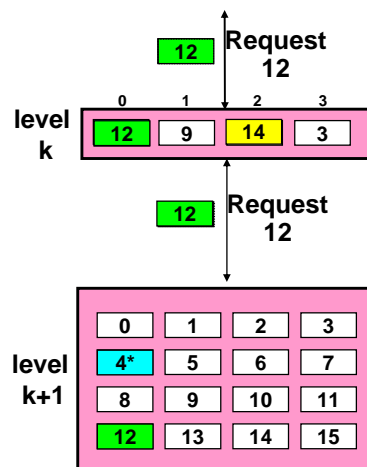


- **Cache:** a smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.
- Fundamental idea of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$.
- Why do memory hierarchies work?
 - Programs tend to access the data at level k more often than they access the data at level $k+1$.
 - Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per bit.

Caching in a memory hierarchy



General caching concepts



- Program needs object d , which is stored in some block b .
- **Cache hit**
 - Program finds b in the cache at level k . E.g., block 14.
- **Cache miss**
 - b is not at level k , so level k cache must fetch it from level $k+1$. E.g., block 12.
 - If level k cache is full, then some current block must be replaced (evicted). Which one is the "victim"?
 - **Placement policy:** where can the new block go? E.g., $b \bmod 4$
 - **Replacement policy:** which block should be evicted? E.g., LRU


Type of cache misses



- **Cold (compulsory) miss:** occurs because the cache is empty.
- **Capacity miss:** occurs when the active cache blocks (working set) is larger than the cache.
- **Conflict miss**
 - Most caches limit blocks at level $k+1$ to a small subset of the block positions at level k , e.g. block i at level $k+1$ must be placed in block $(i \bmod 4)$ at level k .
 - Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block, e.g. Referencing blocks 0, 8, 0, 8, 0, 8, ... would miss every time.

Cache memories

- Cache memories are small, fast SRAM-based memories managed automatically in hardware.
- CPU looks first for data in L1, then in L2, then in main memory.
- Typical system structure:



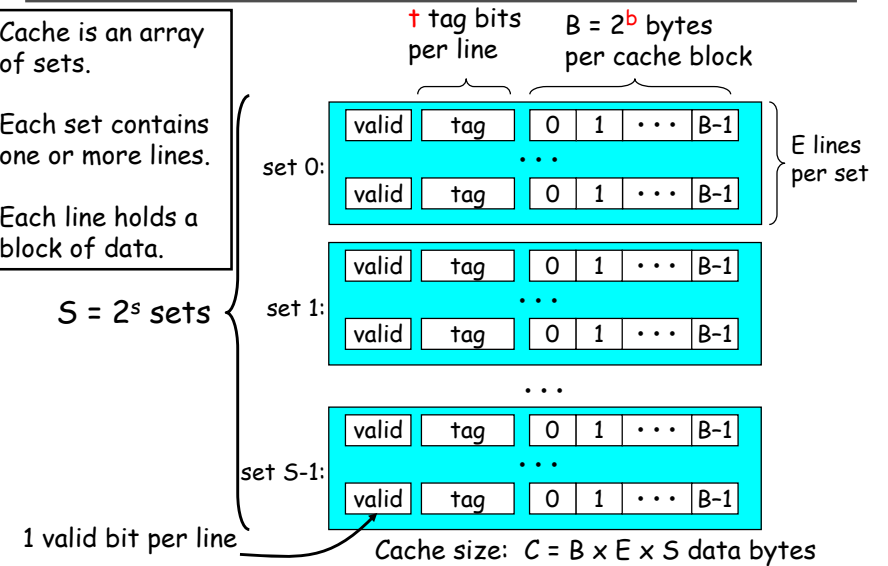
General organization of a cache

Cache is an array of sets.

Each set contains one or more lines.

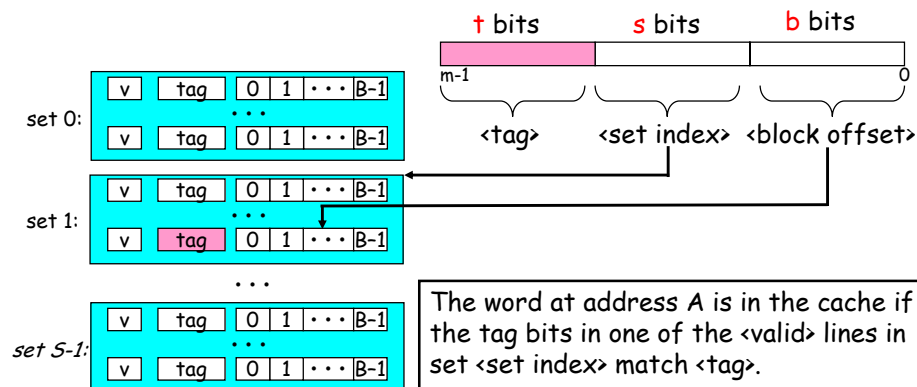
Each line holds a block of data.

$$S = 2^s \text{ sets}$$



Addressing caches

Address A:

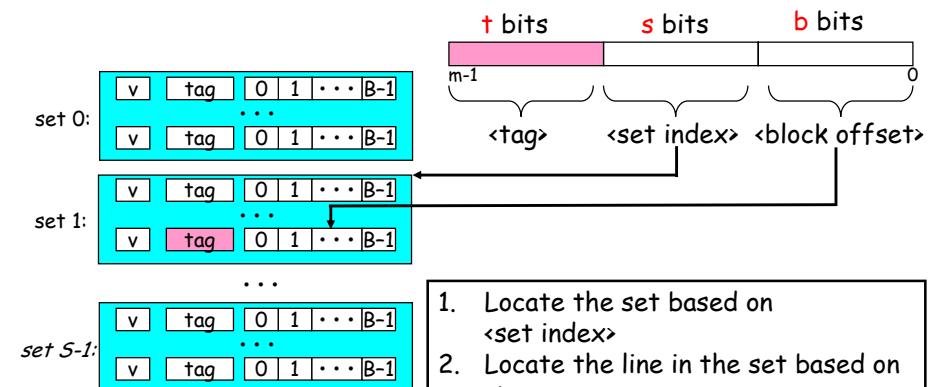


The word at address A is in the cache if the tag bits in one of the $\langle \text{valid} \rangle$ lines in set $\langle \text{set index} \rangle$ match $\langle \text{tag} \rangle$.

The word contents begin at offset $\langle \text{block offset} \rangle$ bytes from the beginning of the block.

Addressing caches

Address A:




1. Locate the set based on $\langle \text{set index} \rangle$
2. Locate the line in the set based on $\langle \text{tag} \rangle$
3. Check that the line is valid
4. Locate the data in the line based on $\langle \text{block offset} \rangle$

Direct-mapped cache



- Simplest kind of cache, easy to build (only 1 tag compare required per access)
- Characterized by exactly one line per set.

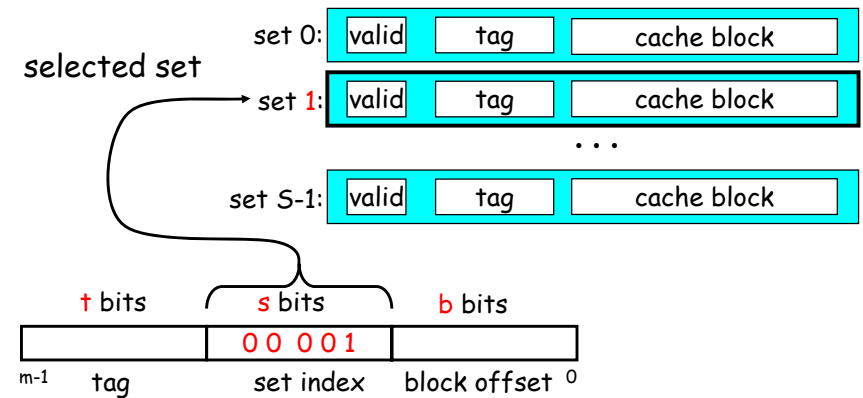


Cache size: $C = B \times S$ data bytes

Accessing direct-mapped caches



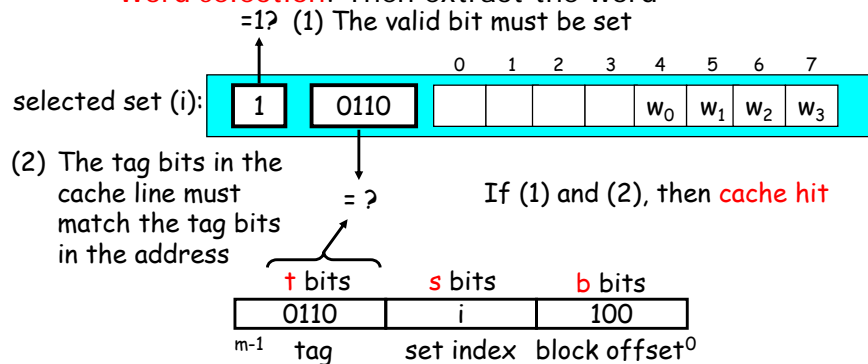
- Set selection
 - Use the set index bits to determine the set of interest.



Accessing direct-mapped caches



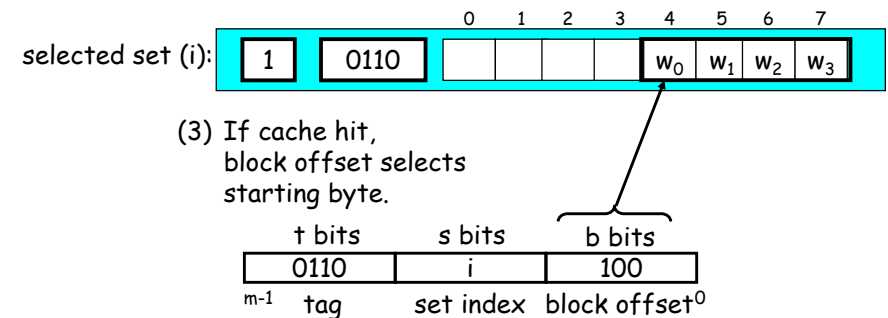
- Line matching and word selection
 - **Line matching**: Find a valid line in the selected set with a matching tag
 - **Word selection**: Then extract the word
 - = 1? (1) The valid bit must be set



Accessing direct-mapped caches



- Line matching and word selection
 - **Line matching**: Find a valid line in the selected set with a matching tag
 - **Word selection**: Then extract the word



Direct-mapped cache simulation



t=1 s=2 b=1

X	XX	X
---	----	---

M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 entry/set

Address trace (reads):

0 [0000₂], miss
1 [0001₂], hit
7 [0111₂], miss
8 [1000₂], miss
0 [0000₂] miss

v tag data

1	0	M[0-1]
1	0	M[6-7]

What's wrong with direct-mapped?



```
float dotprod(float x[8], y[8]) {
    float sum=0.0;
    for (int i=0; i<8; i++)
        sum+= x[i]*y[i];
    return sum;
}
```

block size=16 bytes

element	address	set	element	address	set
x[0]	0	0	y[0]	32	0
x[1]	4	0	y[1]	36	0
x[2]	8	0	y[2]	40	0
x[3]	12	0	y[3]	44	0
x[4]	16	1	y[4]	48	1
x[5]	20	1	y[5]	52	1
x[6]	24	1	y[6]	56	1
x[7]	28	1	y[7]	60	1

Solution? padding




```
float dotprod(float x[12], y[8]) {
    float sum=0.0;
    for (int i=0; i<8; i++)
        sum+= x[i]*y[i];
    return sum;
}
```

element	address	set	element	Address	set
x[0]	0	0	y[0]	48	1
x[1]	4	0	y[1]	52	1
x[2]	8	0	y[2]	56	1
x[3]	12	0	y[3]	60	1
x[4]	16	1	y[4]	64	0
x[5]	20	1	y[5]	68	0
x[6]	24	1	y[6]	72	0
x[7]	28	1	y[7]	76	0

Set associative caches



- Characterized by more than one line per set



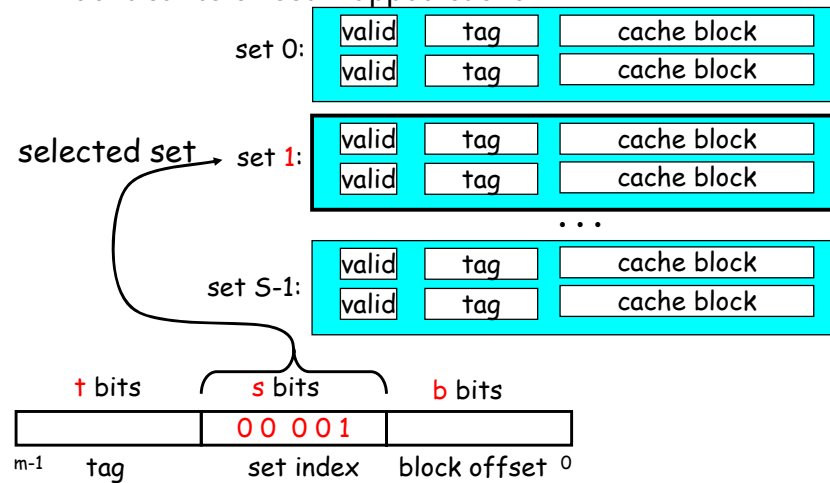
E-way associative cache

Accessing set associative caches



• Set selection

- identical to direct-mapped cache




Accessing set associative caches



• Line matching and word selection

- must compare the tag in each valid line in the selected set.

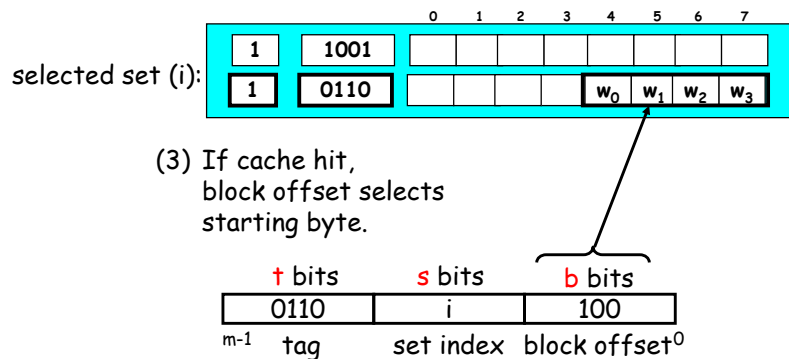


Accessing set associative caches



• Line matching and word selection

- Word selection is the same as in a direct mapped cache



2-Way associative cache simulation



t=2 s=1 b=1
xx x x

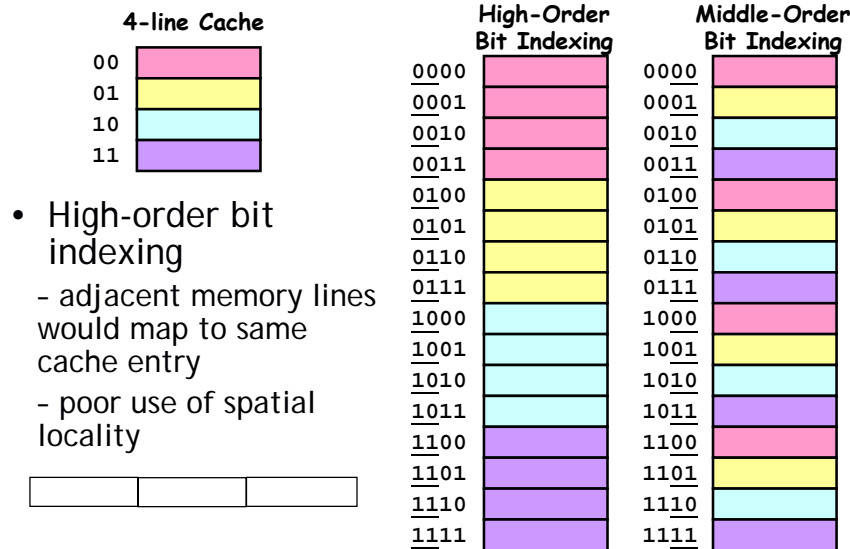
M=16 byte addresses, B=2 bytes/block,
S=2 sets, E=2 entry/set

Address trace (reads):

0 [0000₂], miss
1 [0001₂], hit
7 [0111₂], miss
8 [1000₂], miss
0 [0000₂] hit

v	tag	data
1	00	M[0-1]
1	10	M[8-9]
1	01	M[6-7]
0		

Why use middle bits as index?



What about writes?

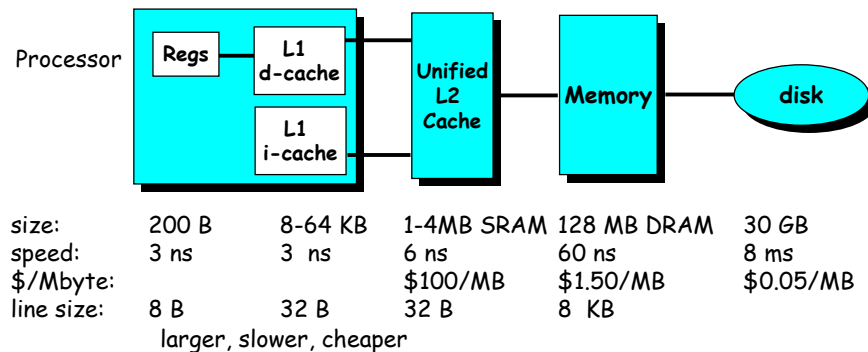


- Multiple copies of data exist:
 - L1
 - L2
 - Main Memory
 - Disk
- What to do when we write?
 - Write-through
 - Write-back (need a dirty bit)
- What to do on a replacement?
 - Depends on whether it is write through or write back

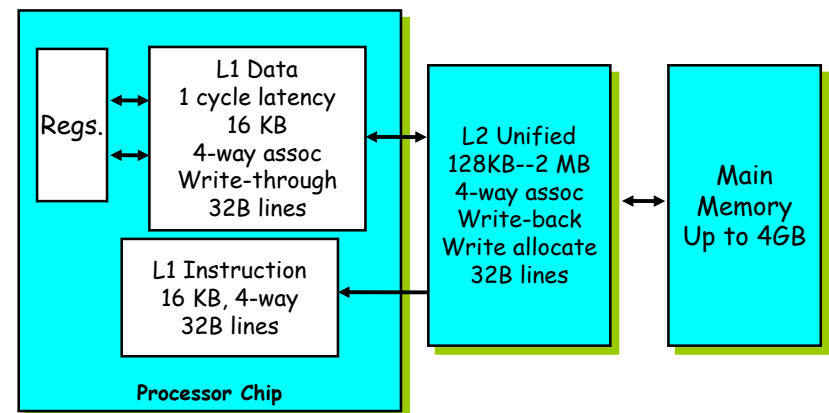
Multi-level caches



- Options: separate **data** and **instruction** caches, or a **unified cache**



Intel Pentium III cache hierarchy



Writing cache friendly code



- Repeated references to variables are good (**temporal locality**)
- Stride-1 reference are good (**spatial locality**)
- Examples: cold cache, 4-byte words, 4-word cache blocks

```
int sum_array_rows(int a[4][8])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = $1/4 = 25\%$

```
int sum_array_cols(int a[4][8])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = **100%**

The memory mountain



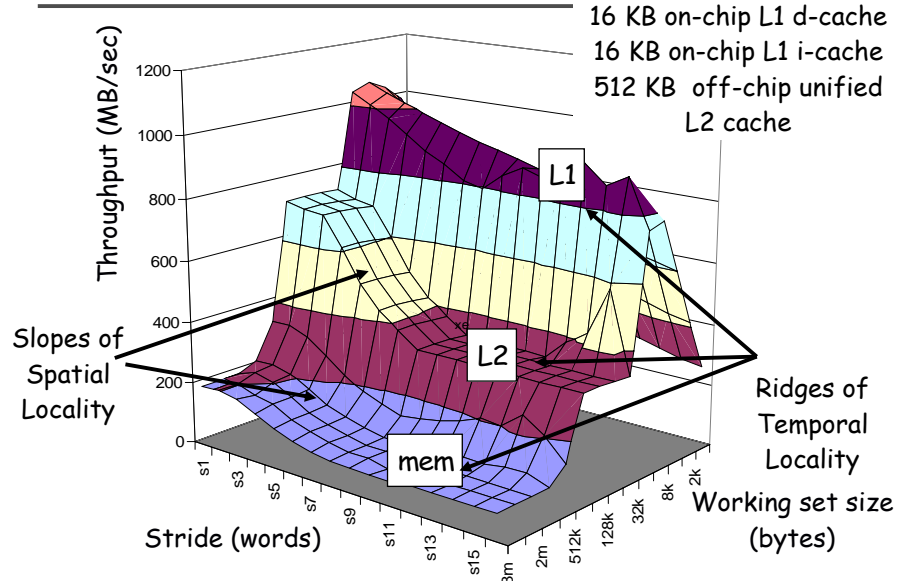
- Read throughput: number of bytes read from memory per second (MB/s)
- Memory mountain
 - Measured read throughput as a function of spatial and temporal locality.
 - Compact way to characterize memory system performance.

```
void test(int elems, int stride) {
    int i, result = 0;
    volatile int sink;
    for (i = 0; i < elems; i += stride)
        result += data[i];
    /* So compiler doesn't optimize away the loop */
    sink = result;
}
```

The memory mountain

Pentium III
550 MHz


16 KB on-chip L1 d-cache
16 KB on-chip L1 i-cache
512 KB off-chip unified L2 cache



Ridges of temporal locality




- Slice through the memory mountain (stride=1)
 - illuminates read throughputs of different caches and memory



A slope of spatial locality



- Slice through memory mountain (size=256KB)
 - shows cache block size.



Matrix multiplication example



- Major cache effects to consider
 - Total cache size
 - Exploit temporal locality and keep the working set small (e.g., use blocking)
 - Block size
 - Exploit spatial locality

```

/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
    
```


Variable sum held in register (pointing to the `sum` variable)

- Description:
 - Multiply N x N matrices
 - $O(N^3)$ total operations
 - Accesses
 - N reads per source element
 - N values summed per destination
 - but may be able to hold in register

Miss rate analysis for matrix multiply



- Assume:
 - Line size = 32B (big enough for four 64-bit words)
 - Matrix dimension (N) is very large
 - Approximate $1/N$ as 0.0
 - Cache is not even big enough to hold multiple rows
- Analysis method:
 - Look at access pattern of inner loop




Matrix multiplication (ijk)



```

/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
    
```




Misses per Inner Loop Iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0

Matrix multiplication (jik)



```
/* jik */
for (j=0; j<n; j++) {
  for (i=0; i<n; i++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum
  }
}
```




Misses per Inner Loop Iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0

Matrix multiplication (kij)



```
/* kij */
for (k=0; k<n; k++) {
  for (i=0; i<n; i++) {
    r = a[i][k];
    for (j=0; j<n; j++)
      c[i][j] += r * b[k][j];
  }
}
```




Misses per Inner Loop Iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25

Matrix multiplication (ikj)



```
/* ikj */
for (i=0; i<n; i++) {
  for (k=0; k<n; k++) {
    r = a[i][k];
    for (j=0; j<n; j++)
      c[i][j] += r * b[k][j];
  }
}
```




Misses per Inner Loop Iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25

Matrix multiplication (jki)



```
/* jki */
for (j=0; j<n; j++) {
  for (k=0; k<n; k++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```



Misses per Inner Loop Iteration:


<u>A</u>	<u>B</u>	<u>C</u>
1.0	0.0	1.0

Matrix multiplication (kji)



```
/* kji */
for (k=0; k<n; k++) {
  for (j=0; j<n; j++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```

Inner loop:



Misses per Inner Loop Iteration:

<u>A</u>	<u>B</u>	<u>C</u>
1.0	0.0	1.0

Summary of matrix multiplication



```
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```

ijk (& jik):

- 2 loads, 0 stores
- misses/iter = **1.25**

```
for (k=0; k<n; k++) {
  for (i=0; i<n; i++) {
    r = a[i][k];
    for (j=0; j<n; j++)
      c[i][j] += r * b[k][j];
  }
}
```

kij (& ikj):

- 2 loads, 1 store
- misses/iter = **0.5**

```
for (j=0; j<n; j++) {
  for (k=0; k<n; k++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```

jki (& kji):


- 2 loads, 1 store
- misses/iter = **2.0**

Pentium matrix multiply performance



- Miss rates are helpful but not perfect predictors.

• Code scheduling matters, too.



Improving temporal locality by blocking



- Example: Blocked matrix multiplication
 - Here, "block" does not mean "cache block".
 - Instead, it mean a sub-block within the matrix.
 - Example: N = 8; sub-block size = 4

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix}$$

Key idea: Sub-blocks (i.e., A_{xy}) can be treated just like scalars.

$$C_{11} = A_{11}B_{11} + A_{12}B_{21} \quad C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21} \quad C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

Blocked matrix multiply (bijk)



```

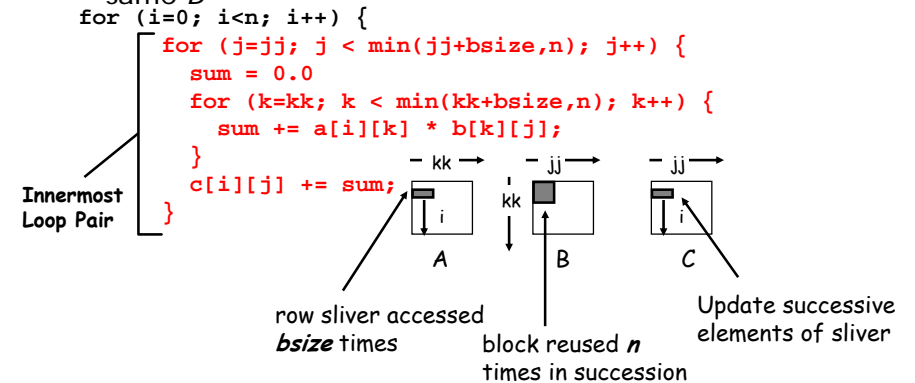
for (jj=0; jj<n; jj+=bsize) {
    for (i=0; i<n; i++)
        for (j=jj; j < min(jj+bsize,n); j++)
            c[i][j] = 0.0;
    for (kk=0; kk<n; kk+=bsize) {
        for (i=0; i<n; i++) {
            for (j=jj; j < min(jj+bsize,n); j++) {
                sum = 0.0
                for (k=kk; k < min(kk+bsize,n); k++) {
                    sum += a[i][k] * b[k][j];
                }
                c[i][j] += sum;
            }
        }
    }
}

```

Blocked matrix multiply analysis



- Innermost loop pair multiplies a $1 \times bsize$ sliver of A by a $bsize \times bsize$ block of B and accumulates into $1 \times bsize$ sliver of C
- Loop over i steps through n row slivers of A & C , using same B




Blocked matrix multiply performance



- Blocking (bijk and bikj) improves performance by a factor of two over unblocked versions (ijk and jik)

- relatively insensitive to array size.



Concluding observations



- Programmer can optimize for cache performance
 - How data structures are organized
 - How data are accessed
 - Nested loop structure
 - Blocking is a general technique
- All systems favor "cache friendly code"
 - Getting absolute optimum performance is very platform specific
 - Cache sizes, line sizes, associativities, etc.
 - Can get most of the advantage with generic code
 - Keep working set reasonably small (temporal locality)
 - Use small strides (spatial locality)