

# MONGODB

# Why NoSQL?



2

- Semantics over the web based on system tags
  - ▣ metadata attached to an part of an web page
- Tags on the Web
  - ▣ ~50 billion pages – horizontal scalability (distribution)
  - ▣ Only basic structure – different content
- A numbers of web information systems
  - ▣ Concurrency access
  - ▣ Replication

# Why MongoDB?

- ❑ Unknown data structure
  - ▣ Horizontal and vertical differences
- ❑ Access by tag ID and ID of tagged document
  - ▣ Indexes for multiple attributes
- ❑ Read operation to whole objects
  - ▣ Each object == One entry in DB
  - ▣ Objects are stored in binary format
- ❑ Parallelized analysis
  - ▣ MapReduce straight over data

# Prominent Users

- AppScale, **bit.ly**, Business Insider, **CERN LHC**, craigslist, diaspora, Disney Interactive Media Group, **EA**, foursquare, **GitHub**, MTV Networks, **SAP**, Shutterfly, **SourceForge**, The Guardian, **The New York Times**, Etsy, Thumbtack, Uber, Wordnik, Springer, Chicago Tribune, Viacom, ...

# Shutterfly: Photo Metadata



5



## □ Problem

- ▣ More flexibility is needed than Oracle delivers

## □ Results

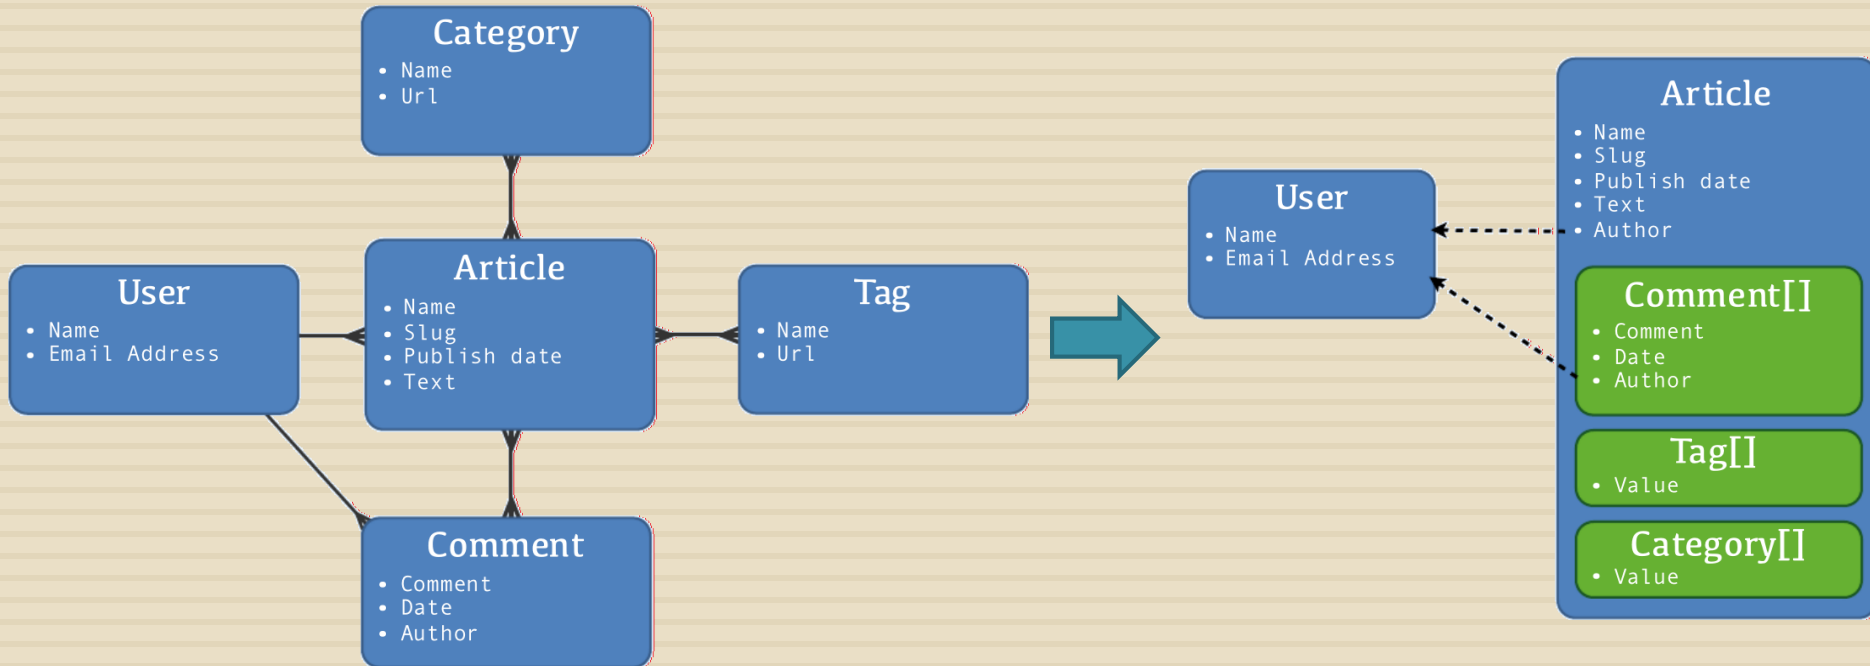
- ▣ 500% cost reduction
- ▣ 900% performance improvement

# Wordnik: Online Dictionary

6

- Problem
  - ▣ MySQL could not scale to handle 5B+ documents
- Results
  - ▣ 20x performance improvement

# Main Idea – Objects in Objects



# Internal Structure

- Databases
- Collections
  - ▣ Equivalent to relational database tables
- Documents
  - ▣ BSON objects
  - ▣ Attributes' data types: object, object id, string, integer, boolean, double, null, array, date, timestamp, binary data, regular expression, code



# BSON (Binary JSON)

## □ JavaScript Object Notation

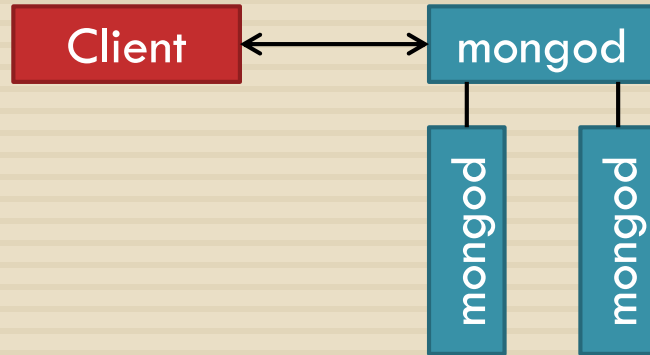
```
{  
  "_id" : CSUUID("56a1f2c8-c1fd"),  
  "UserID" : 777,  
  "MovieID" : 1901,  
  "Value" : 4,  
  "Timestamp" : NumberLong(975604013)  
}
```

# Architecture

- Single node/database

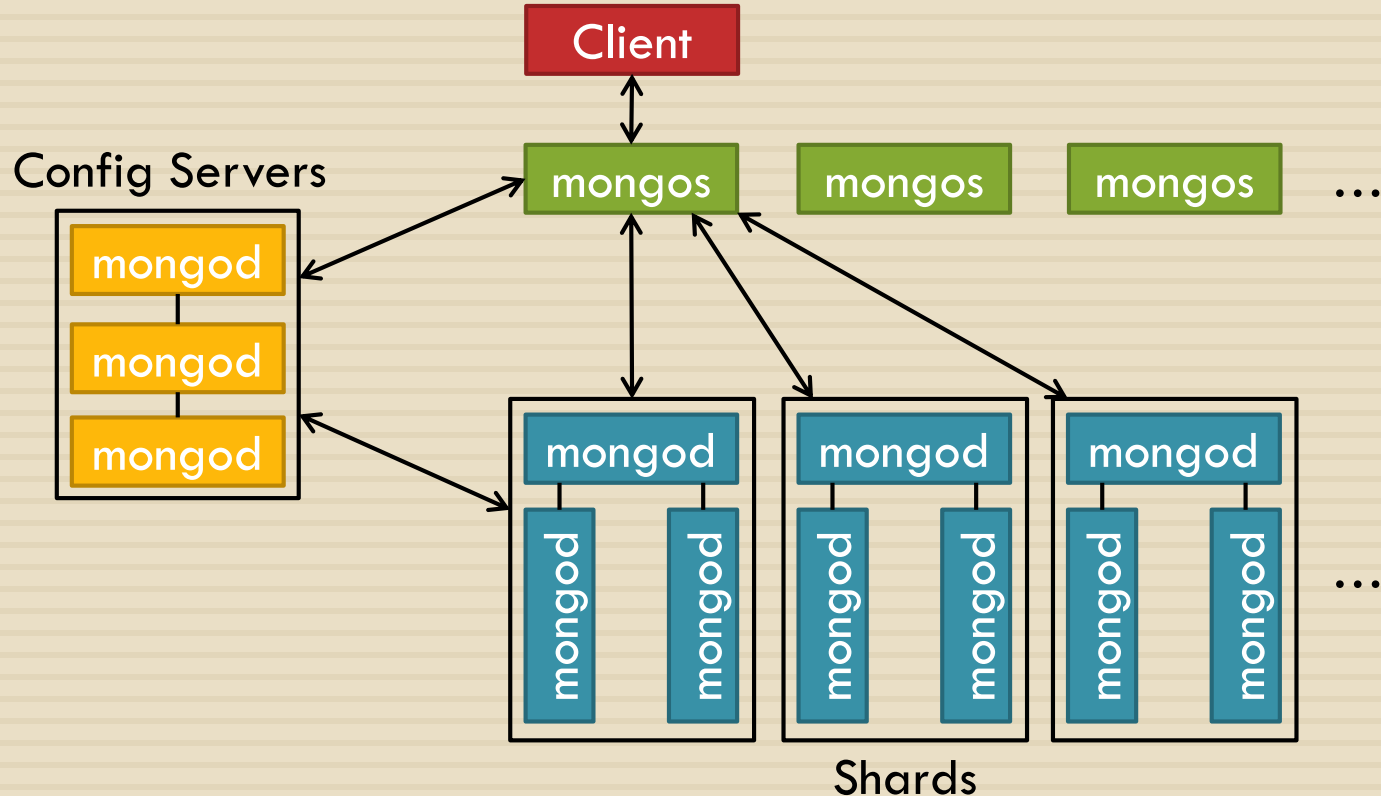


- Single database replicated over several nodes



- Shared (partitioned) database

# Architecture - Partitioned database

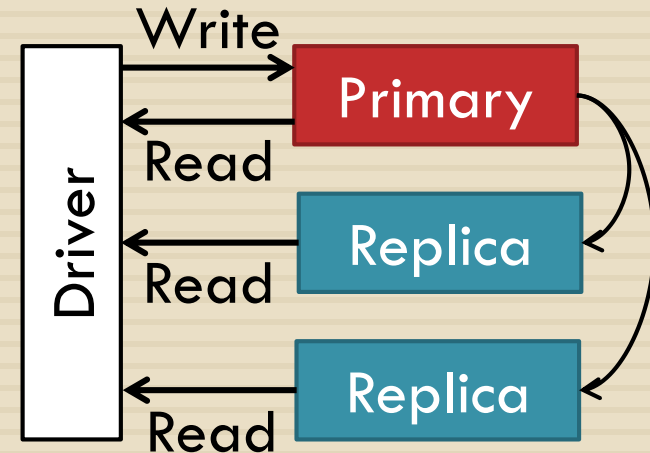


# Replication

- Master slave
  - ▣ 1 mongod is initialized as master – allows writing
  - ▣  $\geq 1$  mongods are initialized as slaves – replicated directly from the master
- Replica set
  - ▣ Type of Master slave
  - ▣ Master (primary node) is elected

# Consistency

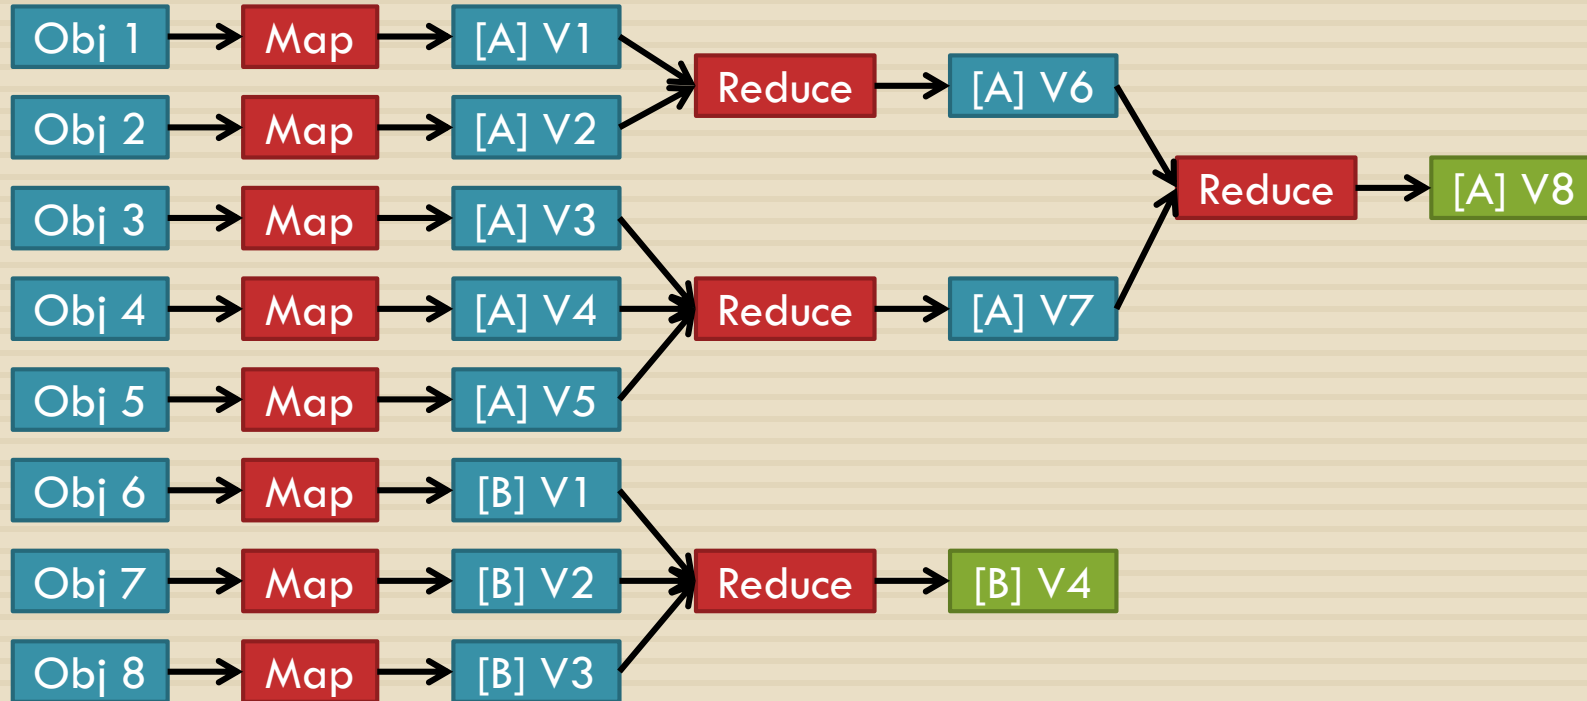
- Managed by drivers
  - Strong consistency
    - ▣ Reading from the primary node
    - ▣ E.g., updating of source code
- Eventual consistency
  - ▣ Reading from any live node
  - ▣ E.g., wiki page



# Querying

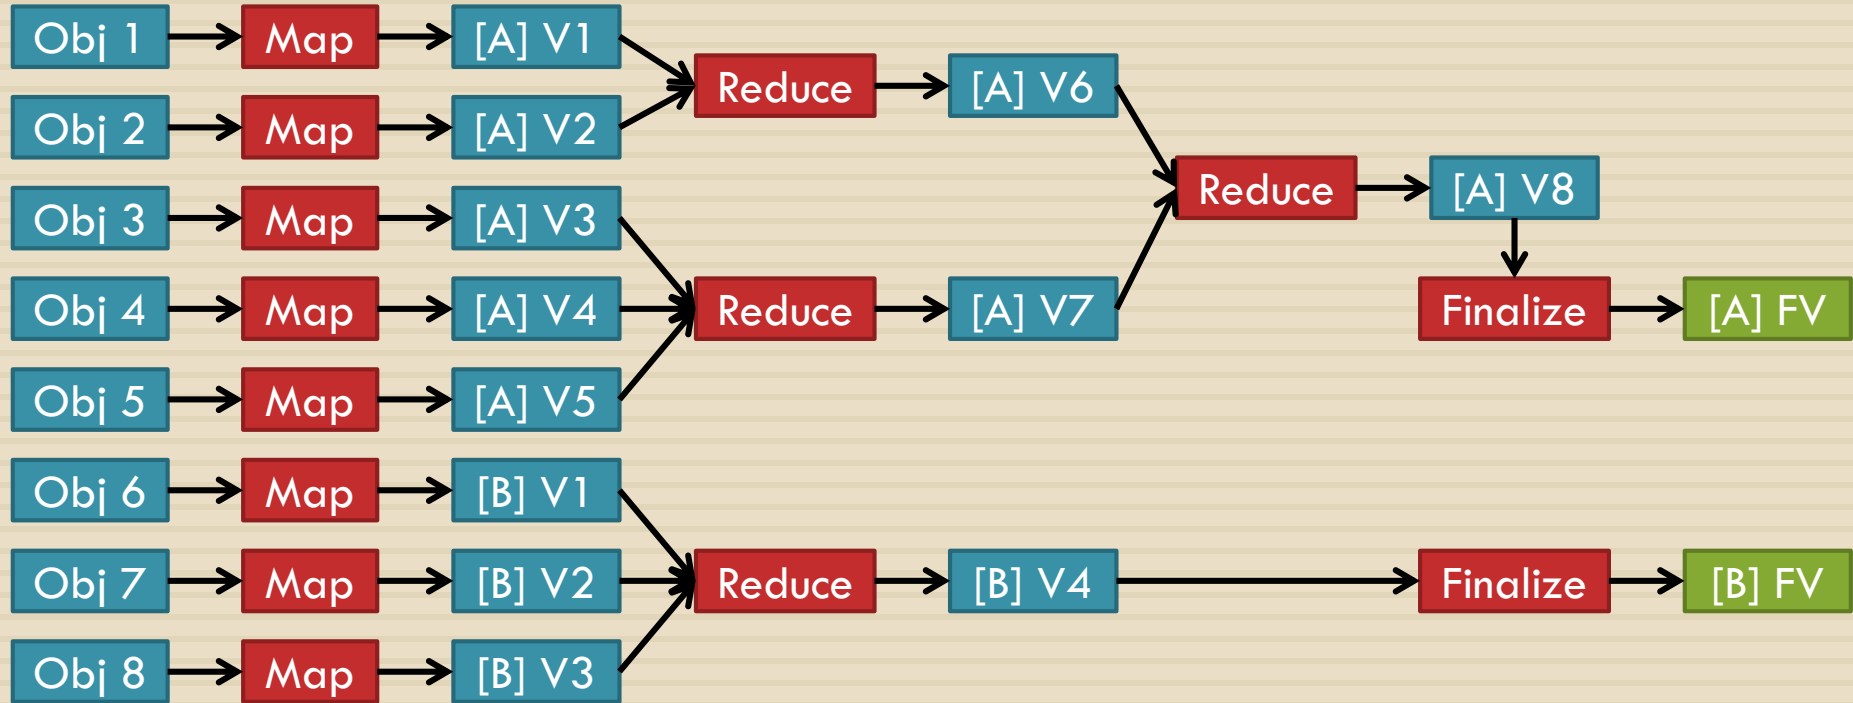
- Indexing
- Querying objects' attributes
  - ▣ Logic expressions, mathematical operations, regexes, ...
  - ▣ Sort, skip, limit
  - ▣ Result: The cursor to the array of queried objects
  - ▣ `db.foo.find( { name : "bob" , $or : [ { a : 1 } , { b : 2 } ] } )`
- MapReduce + Finalize
  - ▣ Result: The view

# MapReduce



# MapReduceFinalize

16





# Updating

- ❑ Several confirmation techniques
- ❑ No transactions
- ❑ Atomic operations
  - ▣ `db.shapes.update( { type: "square" }, { $inc: { x: 5 } } )`
- ❑ Blocking – does not support sharding
  - ▣ `db.students.update({score: {$gt: 60}, $atomic: true}, {$set: {pass: true}}, false, true)`
- ❑ Find and Modify
  - ▣ Returns old/new value

# Connectors

## □ Drivers

- mongodb.org: C, **C++**, Erlang, Haskell, **Java**, Javascript, .NET (**C#**, PowerShell, etc), Perl, **PHP**, Python, **Ruby**, Scala
- Community: ActionScript3, Clojure, ColdFusion, D, Delphi, Entity, Factor, Go, Groovy, Lisp, Lua, MatLab, **node.js**, Objective C, Opa, Prolog, R, REST, Racket, Smalltalk

## □ REST add-on

## □ JavaScript console

# Demonstration

- Dataset: 1 M MovieLens

- ▣ ID, UserID, MovieID,  
Value, TimeStamp

- JavaScript Console

- C# driver

```
> use Test
switched to db Test
> db.Annotations.save({ _id: "http://fiit.sk/Anot2",IsDeleted: false })
> db.Annotations.find().forEach(printjson)
{ "_id" : "http://fiit.sk/Anot1", "IsDeleted" : true }
{ "_id" : "http://fiit.sk/Anot2", "IsDeleted" : false }
> annot = db.Annotations.findOne({_id: "http://fiit.sk/Anot2"})
{ "_id" : "http://fiit.sk/Anot1", "IsDeleted" : false }
> annot.IsDeleted = true
true
> db.Annotations.save(annot)
> db.Annotations.find().forEach(printjson)
{ "_id" : "http://fiit.sk/Anot1", "IsDeleted" : true }
{ "_id" : "http://fiit.sk/Anot2", "IsDeleted" : true }
```

# Resources

- Website:

- <http://www.mongodb.org/>



- SQL to MongoDB

- <http://www.mongodb.org/display/DOCS/SQL+to+Mongo+Mapping+Chart>

- MongoDB&hadoop

- <http://www.slideshare.net/spf13/mongodb-and-hadoop>

- Practical Replication Video

- <http://www.10gen.com/presentations/mongosv-2011/a-mongodb-replication-primer-replica-sets-in-practice>